
Parallel Programming

Distributed-Memory Programming with MPI (Part II)

Professor Yi-Ping You (游逸平)

Department of Computer Science

<http://www.cs.nctu.edu.tw/~ypyou/>



Acknowledgement

- This set of slides is mainly developed by Prof. Thomas Sterling and Dr. Steve Brandt from Louisiana State University



Review of Basic MPI Calls

- In review, the 6 main MPI calls:
 - MPI_Init
 - MPI_Finalize
 - MPI_Comm_size
 - MPI_Comm_rank
 - MPI_Send
 - MPI_Recv
- Include MPI Header file
 - #include "mpi.h"
- Basic MPI Datatypes
 - MPI_INT, MPI_FLOAT,

Collective Calls

- A communication pattern that encompasses all processes within a communicator is known as **collective communication**
- MPI has several collective communication calls, the most frequently used are:
 - Synchronization
 - Barrier
 - Communication
 - Broadcast
 - Gather & Scatter
 - All Gather
 - Reduction
 - Reduce
 - AllReduce

Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary


MPI Collective Calls : Barrier

Function: **MPI_Barrier()**

```
int MPI_Barrier (  
    MPI_Comm comm )
```

Description:

Creates barrier synchronization in a communicator group *comm*. Each process, when reaching the MPI_Barrier call, blocks until all the processes in the group reach the same MPI_Barrier call.



Example: MPI_Barrier()

```
#include <stdio.h>
#include "mpi.h"

int main (int argc, char *argv[]){
    int      rank, size, len;
    char      name[MPI_MAX_PROCESSOR_NAME];
    MPI_Init(&argc, &argv);
    MPI_Barrier(MPI_COMM_WORLD);

    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Get_processor_name(name, &len);
    MPI_Barrier(MPI_COMM_WORLD);

    printf ("Hello world! Process %d of %d on %s\n", rank, size, name);
    MPI_Finalize();
    return 0;
}
```

```
[cdekate@celeritas collective]$ mpirun -np 8 barrier
Hello world! Process 0 of 8 on celeritas.cct.lsu.edu
Writing logfile....
Finished writing logfile.
Hello world! Process 4 of 8 on compute-0-3.local
Hello world! Process 1 of 8 on compute-0-0.local
Hello world! Process 3 of 8 on compute-0-2.local
Hello world! Process 6 of 8 on compute-0-5.local
Hello world! Process 7 of 8 on compute-0-6.local
Hello world! Process 5 of 8 on compute-0-4.local
Hello world! Process 2 of 8 on compute-0-1.local
[cdekate@celeritas collective]$
```



Topics

- MPI Collective Calls: Synchronization Primitives
- **MPI Collective Calls: Communication Primitives**
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

MPI Collective Calls : Broadcast

Function: **MPI_Bcast()**

```
int MPI_Bcast (  
    void          *message,  
    int           count,  
    MPI_Datatype  datatype,  
    int           root,  
    MPI_Comm      comm )
```



Description:

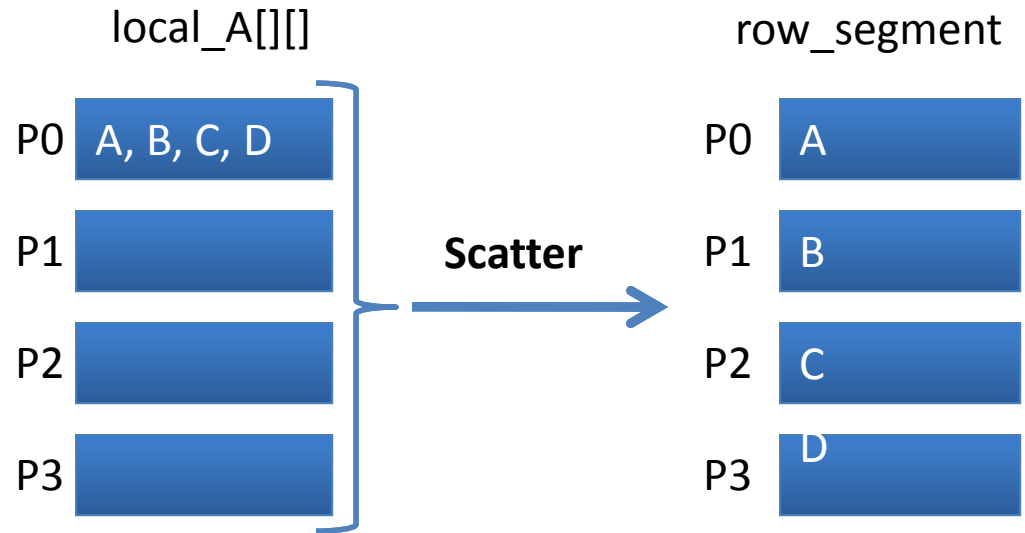
A collective communication call where a single process sends the same data contained in the *message* to every process in the communicator. By default a tree like algorithm is used to broadcast the message to a block of processors, a linear algorithm is then used to broadcast the message from the first process in a block to all other processes. All the processes invoke the *MPI_Bcast* call with the same arguments for *root* and *comm*.

```
float          endpoint[2];  
...  
MPI_Bcast(endpoint, 2, MPI_FLOAT, 0, MPI_COMM_WORLD);
```

MPI Collective Calls : Scatter

Function: **MPI_Scatter()**

```
int MPI_Scatter (  
    void                *sendbuf,  
    int                 send_count,  
    MPI_Datatype         send_type,  
    void                *recvbuf,  
    int                 recv_count,  
    MPI_Datatype         recv_type,  
    int                 root,  
    MPI_Comm             comm)
```



Description:

http://www-unix.mcs.anl.gov/mpi/www/www3/MPI_Scatter.html


MPI_Scatter splits the data referenced by the *sendbuf* on the process with rank *root* into *p* segments each of which consists of *send_count* elements of type *send_type*. The first segment is sent to process0 and the second segment to process1. The send arguments are significant on the process with rank *root*.

```
...  
MPI_Scatter(&(local_A[0][0]), n/p, MPI_FLOAT, row_segment, n/p, MPI_FLOAT, 0, MPI_COMM_WORLD);  
...
```

MPI Collective Calls : Gather

Function: **MPI_Gather()**

```
int MPI_Gather (  
    void                *sendbuf,  
    int                 send_count,  
    MPI_Datatype         sendtype,  
    void                *recvbuf,  
    int                 recvcount,  
    MPI_Datatype         recvtype,  
    int                 root,  
    MPI_Comm             comm )
```



Description: http://www-unix.mcs.anl.gov/mpi/www/www3/MPI_Gather.html

MPI_Gather collects the data referenced by *sendbuf* from each process in the communicator *comm*, and stores the data in process rank order on the process with rank *root* in the location referenced by *recvbuf*. The *recv* parameters are only significant.


```
...  
    MPI_Gather(local_x, n/p, MPI_FLOAT, global_x, n/p, MPI_FLOAT, 0, MPI_COMM_WORLD);  
...
```

MPI Collective Calls : All Gather

Function: **MPI_Allgather()**

```
int MPI_Allgather (  
    void          *sendbuf,  
    int           send_count,  
    MPI_Datatype   sendtype,  
    void          *recvbuf,  
    int           recvcount,  
    MPI_Datatype   recvttype,  
    MPI_Comm       comm )
```

http://www-unix.mcs.anl.gov/mpi/www/www3/MPI_Allgather.html



Description:

MPI_Allgather gathers the content from the send buffer (*sendbuf*) on each process. The effect of this call is similar to executing MPI_Gather() *p* times with a different process acting as the root.

```
for (root=0; root<p; root++)  
    MPI_Gather(local_x, n/p, MPI_FLOAT, global_x, n/p, MPI_FLOAT, root, MPI_COMM_WORLD);  
...
```

CAN BE REPLACED WITH :

```
MPI_Allgather(local_x, local_n, MPI_FLOAT, global_x, local_n, MPI_FLOAT, MPI_COMM_WORLD);
```

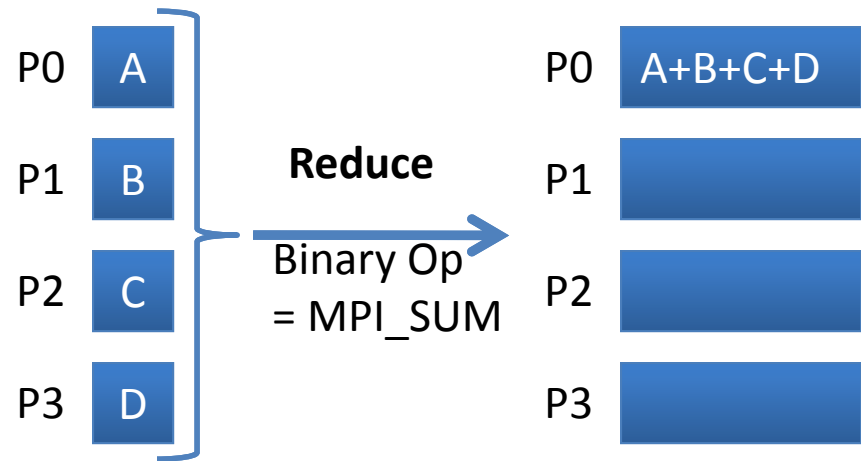
Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- **MPI Collective Calls: Reduction Primitives**
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

MPI Collective Calls : Reduce

Function: **MPI_Reduce()**

```
int MPI_Reduce (  
    void          *operand,  
    void          *result,  
    int           count,  
    MPI_Datatype  datatype,  
    MPI_Op        operator,  
    int           root,  
    MPI_Comm      comm )
```



http://www-unix.mcs.anl.gov/mpi/www/www3/MPI_Reduce.html

Description:

A collective communication call where all the processes in a communicator contribute data that is combined using binary operations (MPI_Op) such as addition, max, min, logical, and, etc. MPI_Reduce combines the operands stored in the memory referenced by *operand* using the operation *operator* and stores the result in **result*. MPI_Reduce is called by all the processes in the communicator *comm* and for each of the processes *count*, *datatype* *operator* and *root* remain the same.

```
...  
MPI_Reduce(&local_integral, &integral, 1, MPI_FLOAT, MPI_SUM, 0, MPI_COMM_WORLD);  
...
```

MPI Binary Operations

- MPI binary operators are used in the MPI_Reduce function call as one of the parameters. MPI_Reduce performs a global reduction operation (dictated by the MPI binary operator parameter) on the supplied operands.
- Some of the common MPI Binary Operators used are :

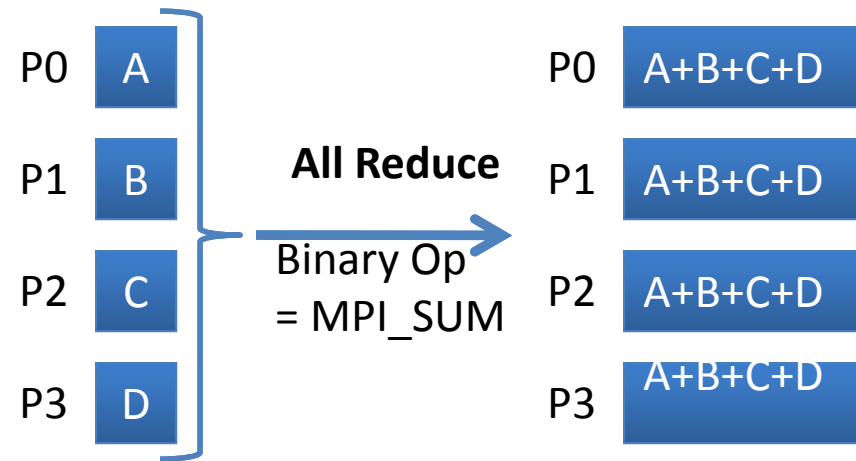
Operation Name Meaning	
MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI_LAND	Logical And
MPI_BAND	Bitwise And
MPI_LOR	Logical Or
MPI_BOR	Bitwise Or
MPI_LXOR	Logical XOR
MPI_BXOR	Bitwise XOR
MPI_MAXLOC	Maximum and location of max.
MPI_MINLOC	Maximum and location of min.

```
MPI_Reduce(&local_integral,  
&integral, 1, MPI_FLOAT,  
MPI_SUM, 0,  
MPI_COMM_WORLD) ;
```


MPI Collective Calls : All Reduce

Function: `MPI_Allreduce()`
on:

```
int MPI_Allreduce (
    void *sendbuf,
    void *recvbuf,
    int count,
    MPI_Datatype datatype,
    MPI_Op op,
    MPI_Comm comm )
```



Description:

MPI_Allreduce is used exactly like MPI_Reduce, except that the result of the reduction is returned on all processes, as a result there is no *root* parameter.

```
...  
MPI_Allreduce(&integral, &integral, 1, MPI_FLOAT, MPI_SUM, MPI_COMM_WORLD);  
...
```

Parallel Trapezoidal Rule

Send, Recv

```
#include <stdio.h>
#include "mpi.h"

main(int argc, char** argv) {
    int    my_rank; /* My process rank */
    int    p;      /* The number of processes */
    float  a = 0.0; /* Left endpoint */
    float  b = 1.0; /* Right endpoint */
    int    n = 1024; /* Number of trapezoids */
    float  h;      /* Trapezoid base length */
    float  local_a; /* Left endpoint my process */
    float  local_b; /* Right endpoint my process */
    int    local_n; /* Number of trapezoids for my
calculation */
    float  integral; /* Integral over my interval */
    float  total;    /* Total integral */
    int    source;   /* Process sending integral */
    int    dest = 0; /* All messages go to 0 */
    int    tag = 0;
    MPI_Status status;
```

```
float Trap(float local_a, float local_b, int local_n,
float h); /* Calculate local integral */

/* Let the system do what it needs to start up MPI */
MPI_Init(&argc, &argv);

/* Get my process rank */
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

/* Find out how many processes are being used */
MPI_Comm_size(MPI_COMM_WORLD, &p);

h = (b-a)/n; /* h is the same for all processes */
local_n = n/p; /* So is the number of trapezoids */

/* Length of each process' interval of
 * integration = local_n*h. So my interval
 * starts at: */
local_a = a + my_rank*local_n*h;
local_b = local_a + local_n*h;
integral = Trap(local_a, local_b, local_n, h);
```

Parallel Trapezoidal Rule

Send, Recv

```
if (my_rank == 0) {
    total = integral;
    for (source = 1; source < p; source++) {
        MPI_Recv(&integral, 1, MPI_FLOAT, source, tag,
                MPI_COMM_WORLD, &status);
        total = total + integral;
    }
} else {
    MPI_Send(&integral, 1, MPI_FLOAT, dest,
            tag, MPI_COMM_WORLD);
}
if (my_rank == 0) {
    printf("With n = %d trapezoids, our estimate\n",
        n);
    printf("of the integral from %f to %f = %f\n",
        a, b, total);
}
MPI_Finalize();
} /* main */
```

Trapezoidal Example Adapted from Parallel Programming in MPI P.Pacheco Ch 4

Parallel Trapezoidal Rule

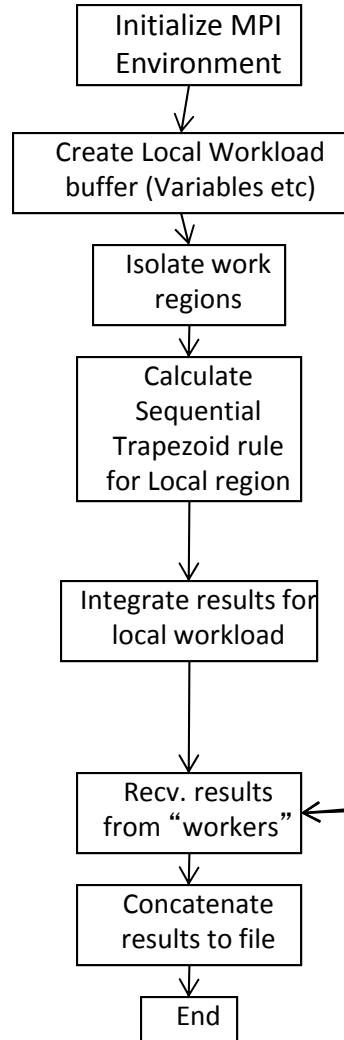
Send, Recv

```
float Trap(  
    float local_a /* in */,  
    float local_b /* in */,  
    int local_n /* in */,  
    float h /* in */) {  
  
    float integral; /* Store result in integral */  
    float x;  
    int i;  
  
    float f(float x); /* function we're integrating */  
  
    integral = (f(local_a) + f(local_b))/2.0;  
    x = local_a;  
    for (i = 1; i <= local_n-1; i++) {  
        x = x + h;  
        integral = integral + f(x);  
    }  
    integral = integral*h;  
    return integral;  
} /* Trap */  
float f(float x) {  
    float return_val;  
    /* Calculate f(x). */  
    /* Store calculation in return_val. */  
    return_val = x*x;  
    return return_val;  
} /* f */
```

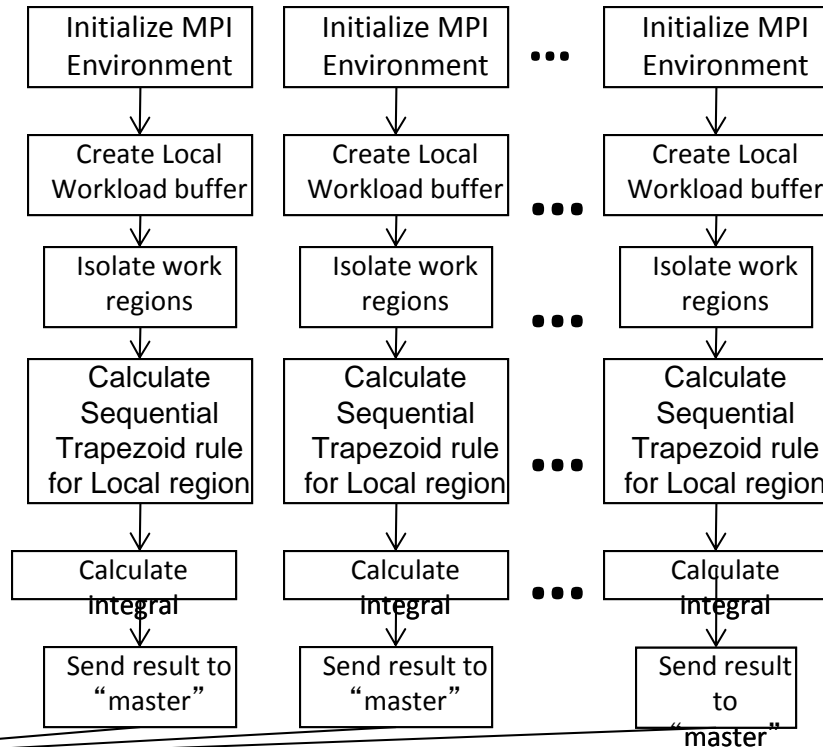
Trapezoidal Example Adapted from Parallel Programming in MPI P.Pacheco Ch 4

Flowchart for Parallel Trapezoidal Rule

MASTER



WORKERS



Trapezoidal Rule : with MPI_Bcast, MPI_Reduce

```
#include <stdio.h>
#include <stdlib.h>
```

```
/* We'll be using MPI routines, definitions, etc. */
#include "mpi.h"
```

```
main(int argc, char** argv) {
    int    my_rank; /* My process rank */
    int    p;      /* The number of processes */
    float  endpoint[2]; /* Left and right */
    int    n = 1024; /* Number of trapezoids */
    float  h;      /* Trapezoid base length */
    float  local_a; /* Left endpoint my process */
    float  local_b; /* Right endpoint my process */
    int    local_n; /* Number of trapezoids for */
                /* my calculation */
    float  integral; /* Integral over my interval */
    float  total;   /* Total integral */
    int    source;  /* Process sending integral */
    int    dest = 0; /* All messages go to 0 */
    int    tag = 0;
    MPI_Status status;
```

```
float Trap(float local_a, float local_b, int local_n,
           float h); /* Calculate local integral */
```

```
MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
if (argc != 3) {
    if (my_rank==0)
        printf("Usage: mpirun -np <numprocs> trapezoid
<left> <right>\n");
    MPI_Finalize();
    exit(0);
}
```

```
if (my_rank==0) {
    endpoint[0] = atof(argv[1]); /* left endpoint */
    endpoint[1] = atof(argv[2]); /* right endpoint */
}
```

```
MPI_Bcast(endpoint, 2, MPI_FLOAT,
0, MPI_COMM_WORLD);
```

Trapezoidal Rule : with MPI_Bcast, MPI_Reduce

```
h = (endpoint[1]-endpoint[0])/n; /* h is the same for all processes
*/
local_n = n/p; /* so is the number of trapezoids */
if (my_rank == 0) printf("a=%f, b=%f, Local number of
trapezoids=%d\n", endpoint[0], endpoint[1], local_n);

local_a = endpoint[0] + my_rank*local_n*h;
local_b = local_a + local_n*h;
integral = Trap(local_a, local_b, local_n, h);
```

**MPI_Reduce(&integral, &total, 1,
MPI_FLOAT, MPI_SUM, 0,
MPI_COMM_WORLD);**

```
if (my_rank == 0) {
    printf("With n = %d trapezoids, our estimate\n",
        n);
    printf("of the integral from %f to %f = %f\n",
        endpoint[0], endpoint[1], total);
}
```

```
MPI_Finalize();
} /* main */
```

```
float Trap(
    float local_a /* in */,
    float local_b /* in */,
    int local_n /* in */,
    float h /* in */) {

    float integral; /* Store result in integral */
    float x;
    int i;

    float f(float x); /* function we're integrating */

    integral = (f(local_a) + f(local_b))/2.0;
    x = local_a;
    for (i = 1; i <= local_n-1; i++) {
        x = x + h;
        integral = integral + f(x);
    }
    integral = integral*h;
    return integral;
} /* Trap */

float f(float x) {
    float return_val;
    /* Calculate f(x). */
    /* Store calculation in return_val. */
    return_val = x*x;
    return return_val;
} /* f */
```

Trapezoidal Rule : with MPI_Bcast, MPI_Reduce

```
#!/bin/bash
#PBS -N name
#PBS -l walltime=120:00:00,nodes=2:ppn=4
cd /home/lsu00/Demos/l9/trapBcast
pwd
date
PROCS=`wc -l < $PBS_NODEFILE`
mpdboot --file=$PBS_NODEFILE
/usr/lib64/mpich2/bin/mpirun -n $PROCS ./trapBcast 2 25 >>out.txt
mpdallexit
date
```

```
[lsu00@master trapBcast]$ qsub try.pbs
518.master.arete.cct.lsu.edu
[lsu00@master trapBcast]$ cat out.txt
a=2.000000, b=25.000000, Local number of trapezoids=128
With n = 1024 trapezoids, our estimate
of the integral from 2.000000 to 25.000000 = 5205.667969
```


Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- **Derived Datatypes: Introduction**
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

Consolidating Data

- In virtually all distributed-memory systems, communication can be much more expensive than local computation
- We would expect the following pair of **for** loops to be much slower than the single send/receive pair

```
double x[1000];  
.  
.  
.  
if (my_rank == 0)  
    for (i = 0; i < 1000; i++)  
        MPI_Send(&x[i], 1, MPI_DOUBLE, 1, 0, comm);  
else /* my_rank == 1 */  
    for (i = 0; i < 1000; i++)  
        MPI_Recv(&x[i], 1, MPI_DOUBLE, 0, 0, comm, &status);  
  
if (my_rank == 0)  
    MPI_Send(x, 1000, MPI_DOUBLE, 1, 0, comm);  
else /* my_rank == 1 */  
    MPI_Recv(x, 1000, MPI_DOUBLE, 0, 0, comm, &status);
```

Consolidating Data (Cont'd)

- MPI provides three basic approaches to consolidating data that might otherwise require multiple messages:
 - the count argument to the various communication functions,
 - derived datatypes, and
 - MPI_Pack/Unpack (see Exercise 3.20)

Constructing Datatypes

- Creating data structures in C:

```
typedef struct {  
    .  
    .  
    .  
} STRUCT_NAME
```

- For example : In the numerical integration by trapezoidal rule we could create a data structure for storing the attributes of the problem as follows:

```
typedef struct {  
    float a,  
    float b,  
    int n;  
} DATA_INTEGRAL;  
.  
.  
.  
DATA_INTEGRAL intg_data;
```

- What would happen when you use:

```
MPI_Bcast( &intg_data, 1, DATA_INTEGRAL, 0, MPI_COMM_WORLD);
```

ERROR!!!

Intg_data is of the type
DATA_INTEGRAL
NOT an MPI_Datatype

Constructing MPI Datatypes

- MPI allows users to define derived MPI datatypes, using basic datatypes that build during execution time
- These derived data types can be used in the MPI communication calls, instead of the basic predefined datatypes.
- A sending process can pack noncontiguous data into contiguous buffer and send the buffered data to a receiving process that can unpack the contiguous buffer and store the data to noncontiguous location.
- A derived datatype is an opaque object that specifies :
 - A sequence of primitive datatypes
 - A sequence of integer (byte) displacements
- MPI has several functions for constructing derived datatypes :
 - Contiguous
 - Vector
 - Indexed
 - Struct

MPI : Basic Data Types (Review)

MPI datatype	C datatype
MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_BYTE	
MPI_PACKED	

You can also define your own (derived datatypes), such as an array of ints of size 100, or more complex examples, such as a struct or an array of structs

Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- **Derived Datatypes: Contiguous**
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

Derived Datatypes : Contiguous

Function: `MPI_Type_contiguous()`

```
int MPI_Type_contiguous(  
    int count,  
    MPI_Datatype old_type,  
    MPI_Datatype *new_type)
```

Description:

This is the simplest constructor in the MPI derived datatypes. Contiguous datatype constructors create a new datatype by making *count* copies of existing data type (*old_type*)

```
MPI_Datatype rowtype;  
...  
MPI_Type_contiguous(SIZE, MPI_FLOAT, &rowtype);  
MPI_Type_commit(&rowtype);  
...
```

http://www-unix.mcs.anl.gov/mpi/www/www3/MPI_Type_contiguous.html

Example : Derived Datatypes - Contiguous

```
#include "mpi.h"
#include <stdio.h>
#define SIZE 4

int main(argc,argv)
int argc;
char *argv[]; {
int numtasks, rank, source=0, dest, tag=1, i;
MPI_Request req;
float a[SIZE][SIZE] =
    {1.0, 2.0, 3.0, 4.0,
     5.0, 6.0, 7.0, 8.0,
     9.0, 10.0, 11.0, 12.0,
     13.0, 14.0, 15.0, 16.0};
float b[SIZE];

MPI_Status stat;
MPI_Datatype rowtype;

MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
MPI_Type_contiguous(SIZE, MPI_FLOAT,
&rowtype);
MPI_Type_commit(&rowtype);
if (numtasks == SIZE) {
    if (rank == 0) {
        for (i=0; i<numtasks; i++){
            dest = i;
            MPI_Isend(&a[i][0], 1, rowtype, dest,
tag, MPI_COMM_WORLD, &req);
        }
        MPI_Recv(b, SIZE, MPI_FLOAT, source, tag,
MPI_COMM_WORLD, &stat);
        printf("rank= %d  b= %3.1f %3.1f %3.1f %3.1f\n",
rank,b[0],b[1],b[2],b[3]);
    }
    else
        printf("Must specify %d processors. Terminating.\n",SIZE);
MPI_Type_free(&rowtype);
MPI_Finalize();
}
```

<https://computing.llnl.gov/tutorials/mpi/>

Example : Derived Datatypes - Contiguous

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

Declares a 4x4 array of
datatype float

```
#include "mpi.h"
#include <stdio.h>
#define SIZE 4

int main(argc,
int argc;
char *argv[];
int numtasks,
MPI_Request req;
```

```
float a[SIZE][SIZE] =
{1.0, 2.0, 3.0, 4.0,
 5.0, 6.0, 7.0, 8.0,
 9.0, 10.0, 11.0, 12.0,
13.0, 14.0, 15.0, 16.0};
float b[SIZE];
```

```
MPI_Status stat;
MPI_Datatype rowtype;
```

```
MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
MPI_Type_contiguous(SIZE, MPI_FLOAT,
&rowtype);
```

```
MPI_Type_commit(&rowtype);
```

```
if (numtasks == SIZE)
if (rank == 0) {
for (i=0; i<SIZE; i++)
dest = i;
```

```
MPI_Isend(a[i], 1, rowtype,
tag, MPI_COMM_WORLD,
0);
}
```

```
MPI_Recv(b, 1, rowtype,
MPI_COMM_WORLD,
tag,
MPI_STATUS_IGNORE,
&stat);
printf("rank= %d\n",
rank,b[0],b[1],b[2],b[3]);
}
```

```
else
```

```
printf("Must specify %d processors. Terminating.\n",SIZE);
MPI_Type_free(&rowtype);
MPI_Finalize();
}
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

Homogenous datastructure
of size 4 (Type : rowtype)

<https://computing.llnl.gov/tutorials/mpi/>

Example : Derived Datatypes - Contiguous

MPI_Type_contiguous

```
count = 4;  
MPI_Type_contiguous(count, MPI_FLOAT, &rowtype);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

a[4][4]

```
[lsu00@master derived]$ qsub contiguous.pbs  
531.master.arete.cct.lsu.edu  
[lsu00@master derived]$ cat out.txt  
rank= 0  b= 1.0 2.0 3.0 4.0  
rank= 1  b= 5.0 6.0 7.0 8.0  
rank= 3  b= 13.0 14.0 15.0 16.0  
rank= 2  b= 9.0 10.0 11.0 12.0
```

```
MPI_Send(&a[2][0], 1, rowtype, dest, tag, comm);
```

9.0	10.0	11.0	12.0
-----	------	------	------

1 element of
rowtype

<https://computing.llnl.gov/tutorials/mpi/>


Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- **Derived Datatypes: Vector**
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

Derived Datatypes : Vector

Function: `MPI_Type_vector()`

```
int MPI_Type_vector(  
    int count,  
    int blocklen,  
    int stride,  
    MPI_Datatype old_type,  
    MPI_Datatype *newtype )
```



count = 4
blocklen = 2
stride = 4

Description:

Returns a new datatype that represents equally spaced blocks. The spacing between the start of each block is given in units of extent (oldtype). The *count* represents the number of blocks, *blocklen* details the number of elements in each block, *stride* represents the number of elements between start of each block of the *old_type*. The new datatype is stored in *new_type*

```
...  
MPI_Type_vector(SIZE, 1, SIZE, MPI_FLOAT, &columnntype);  
...
```

http://www-unix.mcs.anl.gov/mpi/www/www3/MPI_Type_vector.html

Example : Derived Datatypes - Vector

```
#include "mpi.h"
#include <stdio.h>
#define SIZE 4

int main(argc,argv)
int argc;
char *argv[]; {
int numtasks, rank, source=0, dest, tag=1, i;
MPI_Request req;
float a[SIZE][SIZE] =
    {1.0, 2.0, 3.0, 4.0,
     5.0, 6.0, 7.0, 8.0,
     9.0, 10.0, 11.0, 12.0,
     13.0, 14.0, 15.0, 16.0};
float b[SIZE];

MPI_Status stat;
MPI_Datatype columntype;
```

```
MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
```

```
MPI_Type_vector(SIZE, 1, SIZE, MPI_FLOAT, &columntype);
MPI_Type_commit(&columntype);
```

```
if (numtasks == SIZE) {
    if (rank == 0) {
        for (i=0; i<numtasks; i++)
            MPI_Isend(&a[0][i], 1, columntype, i, tag,
MPI_COMM_WORLD, &req);
    }
}
```

```
MPI_Recv(b, SIZE, MPI_FLOAT, source, tag,
MPI_COMM_WORLD, &stat);
printf("rank= %d  b= %3.1f %3.1f %3.1f %3.1f\n",
    rank,b[0],b[1],b[2],b[3]);
}
else
    printf("Must specify %d processors. Terminating.\n",SIZE);
```

```
MPI_Type_free(&columntype);
MPI_Finalize();
}
```

<https://computing.llnl.gov/tutorials/mpi/>

Example - Derived Datatypes - Vector

```
#include "mpi.h"
#include <stdio.h>
#define SIZE 4
```

```
int main(argc,argv)
int argc;
char *argv[]; {
int numtasks, rank,
MPI_Request req;
```

```
float a[SIZE][SIZE] =
{1.0, 2.0, 3.0, 4.0,
5.0, 6.0, 7.0, 8.0,
9.0, 10.0, 11.0, 12.0,
13.0, 14.0, 15.0, 16.0};
float b[SIZE];
```

```
MPI_Status stat;
MPI_Datatype columntype;
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

Declares a 4x4 array of
datatype float

```
argc,&argv);
rank(MPI_COMM_WORLD, &rank);
size(MPI_COMM_WORLD, &numtasks);
```

```
MPI_Type_create_vector(SIZE, 1, SIZE, MPI_FLOAT, &columntype);
MPI_Type_commit(&columntype);
```

```
for(i=0; i<numtasks; i++) {
for(j=0; j<SIZE; j++) {
```

```
MPI_Isend(&a[0][i], 1, columntype,
MPI_COMM_WORLD, j, &stat);
```

```
MPI_Recv(b, SIZE, MPI_FLOAT,
MPI_COMM_WORLD, j, &stat);
printf("rank= %d b= %3.1f %3.1f %3.1f %3.1f\n",
rank,b[0],b[1],b[2],b[3]);
```

```
}
else
printf("Must specify %d processors. Terminating.\n",SIZE);
```

```
MPI_Type_free(&columntype);
MPI_Finalize();
}
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

Homogenous
datastructure of size 4
(Type : *columntype*)

Example : Derived Datatypes - Vector

MPI_Type_vector

```
count = 4; blocklength = 1; stride = 4;  
MPI_Type_vector(count, blocklength, stride, MPI_FLOAT,  
                &column_type);
```

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

a[4][4]

```
[lsu00@master derived]$ qsub vector.pbs  
562.master.arete.cct.lsu.edu  
[lsu00@master derived]$ cat out.txt  
rank= 0  b= 1.0 5.0 9.0 13.0  
rank= 2  b= 3.0 7.0 11.0 15.0  
rank= 1  b= 2.0 6.0 10.0 14.0  
rank= 3  b= 4.0 8.0 12.0 16.0
```

```
MPI_Send(&a[0][1], 1, column_type, dest, tag, comm);
```

2.0	6.0	10.0	14.0
-----	-----	------	------

1 element of
column_type

<https://computing.llnl.gov/tutorials/mpi/>


Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- **Derived Datatypes: Indexed**
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

Derived Datatypes : Indexed

Function: `MPI_Type_indexed()`

```
int MPI_Type_indexed(  
    int count,  
    int *array_of_blocklengths,  
    int *array_of_displacements,  
    MPI_Datatype oldtype,  
    MPI_datatype *newtype);
```



count = 2
array_of_blocklengths[2] = {2, 4}
array_of_displacements[2] = {0, 4}

Description:

Returns a new datatype that represents count blocks. Each block is defined by an entry in *array_of_blocklengths* and *array_of_displacements*. Displacements are expressed in units of extent(oldtype). The *count* is the number of blocks and the number of entries in *array_of_displacements* (displacement of each block in units of the *oldtype*) and *array_of_blocklengths* (number of instances of *oldtype* in each block).

```
...  
MPI_Type_indexed(2, blocklengths, displacements, MPI_FLOAT,  
&indextype);  
...
```

https://computing.llnl.gov/tutorials/mpi/man/MPI_Type_indexed.txt

Example : Derived Datatypes - Indexed

```
#include "mpi.h"
#include <stdio.h>
#define NELEMENTS 6

int main(argc,argv)
int argc;
char *argv[]; {
int numtasks, rank, source=0, dest, tag=1, i;
MPI_Request req;
int blocklengths[2], displacements[2];
float a[16] =
    {1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0,
      9.0, 10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0};
float b[NELEMENTS];

MPI_Status stat;
MPI_Datatype indextype;

MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
```

```
blocklengths[0] = 4;
blocklengths[1] = 2;
displacements[0] = 5;
displacements[1] = 12;
```

```
MPI_Type_indexed(2, blocklengths,
displacements, MPI_FLOAT, &indextype);
MPI_Type_commit(&indextype);
```

```
if (rank == 0) {
    for (i=0; i<numtasks; i++)
        MPI_Isend(a, 1, indextype, i, tag,
MPI_COMM_WORLD, &req);
}
```

```
MPI_Recv(b, NELEMENTS, MPI_FLOAT, source,
tag, MPI_COMM_WORLD, &stat);
printf("rank= %d  b= %3.1f %3.1f %3.1f %3.1f %3.1f %3.1f\n",
    rank,b[0],b[1],b[2],b[3],b[4],b[5]);
```

```
MPI_Type_free(&indextype);
MPI_Finalize();
}
```

Example : Derived Datatypes - Indexed

```
#include "mpi.h"
#include <stdio.h>
#define NELEMENTS 6
```

```
blocklengths[0] = 4;
blocklengths[1] = 2;
displacements[0] = 5;
displacements[1] = 12;
```

1.0 2.0 3.0 4.0 5.0 6.0 7.0 8.0 9.0 10.0 11.0 12.0 13.0 14.0 15.0 16.0

Declares a[16] array of type float

```
MPI_Request req;
int blocklengths[2], displacements[2];
```

```
float a[16] =
    {1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0,
     9.0, 10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0};
float b[NELEMENTS];
```

```
MPI_Status stat;
MPI_Datatype indextype;
```

```
MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
```

1.0 2.0 3.0 4.0 5.0 6.0 7.0 8.0 9.0 10.0 11.0 12.0 13.0 14.0 15.0 16.0

Creates a new datatype *indextype*

```
MPI_Isend(a, 1, indextype, 1, tag,
MPI_COMM_WORLD, &req);
}
```

```
MPI_Recv(b, NELEMENTS, MPI_FLOAT, source,
tag, MPI_COMM_WORLD, &stat);
```

```
printf("rank= %d b= %3.1f %3.1f %3.1f %3.1f %3.1f %3.1f\n",
rank,b[0],b[1],b[2],b[3],b[4],b[5]);
```


```
MPI_Type_free(&indextype);
MPI_Finalize();
}
```

<https://computing.llnl.gov/tutorials/mpi/>

Example : Derived Datatypes - Indexed

MPI_Type_indexed


count = 2; blocklengths[0] = 4;
 displacements[0] = 5; blocklengths[1] = 2;
 displacements[1] = 12;



```
[lsu00@master derived]$ qsub indexed.pbs
560.master.arete.cct.lsu.edu
[lsu00@master derived]$ cat out.txt
rank= 0  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 2  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 1  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 3  b= 6.0 7.0 8.0 9.0 13.0 14.0
```

```
MPI_Type_indexed(count, blocklengths, displacements, MPI_FLOAT, &indextype);
```

```
MPI_Send(&a, 1, indextype, dest, tag, comm);
```



1 element of
indextype

<https://computing.llnl.gov/tutorials/mpi/>

Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- **Derived Datatypes: Struct**
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

Derived Datatypes : struct

Function: `MPI_Type_struct()`

```
int MPI_Type_struct(  
    int count,  
    int *array_of_blocklengths,  
    MPI_Aint *array_of_displacements,  
    MPI_Datatype *array_of_types,  
    MPI_datatype *newtype);
```

Deprecated since MPI-2.0
Removed since MPI-3.0
Replacement: `MPI_Type_create_struct`

Description:

Returns a new datatype that represents count blocks. Each is defined by an entry in `array_of_blocklengths`, `array_of_displacements` and `array_of_types`. Displacements are expressed in bytes. *count* is an integer that specifies the number of blocks (number of entries in arrays). The *array_of_blocklengths* is the number of elements in each blocks & *array_of_displacements* specifies the byte displacement of each block. The *array_of_types* parameter comprising each block is made of concatenation of type *array_of_types*.

```
...  
MPI_Type_struct(2, blockcounts, offsets, oldtypes, &particletype);  
...
```

https://computing.llnl.gov/tutorials/mpi/man/MPI_Type_struct.txt

Example : Derived Datatype - struct

```
#include "mpi.h"
#include <stdio.h>
#define NELEM 25
int main(argc,argv)
int argc;
char *argv[]; {
int numtasks, rank, source=0, dest, tag=1, i;
typedef struct {
    float x, y, z;
    float velocity;
    int n, type;
} Particle;
Particle p[NELEM], particles[NELEM];
MPI_Datatype particletype, oldtypes[2];
int blockcounts[2];
MPI_Aint offsets[2];

/* MPI_Aint type used to be consistent with syntax of */
/* MPI_Type_get_extent routine */
MPI_Aint lb, float_extent;

MPI_Status stat;

MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
```

```
MPI_Type_get_extent(MPI_FLOAT, &lb, &float_extent);
```

```
blockcounts[0] = 4; blockcounts[1] = 2;
offsets[0] = 0; offsets[1] = 4 * float_extent;
oldtypes[0] = MPI_FLOAT; oldtypes[1] = MPI_INT;
MPI_Type_create_struct(2, blockcounts, offsets, oldtypes,
&particletype);
MPI_Type_commit(&particletype);
```

```
if (rank == 0) {
    for (i=0; i<NELEM; i++) {
        particles[i].x = i * 1.0;
        particles[i].y = i * -1.0;
        particles[i].z = i * 1.0;
        particles[i].velocity = 0.25;
        particles[i].n = i;
        particles[i].type = i % 2;
    }
    for (i=0; i<numtasks; i++)
        MPI_Send(particles, NELEM, particletype, i, tag, MPI_COMM_WORLD);
}
MPI_Recv(p, NELEM, particletype, source, tag, MPI_COMM_WORLD, &stat);
printf("rank= %d  %3.2f %3.2f %3.2f %3.2f %d %d\n", rank,p[3].x,
        p[3].y,p[3].z,p[3].velocity,p[3].n,p[3].type);

MPI_Type_free(&particletype);
MPI_Finalize();
}
```

<https://computing.llnl.gov/tutorials/mpi/>

Example : Derived Datatype - struct

```
#include "mpi.h"
#include <stdio.h>
#define NELEM 25
int main(argc,argv)
int argc;
char *argv[]; {
int numtasks, rank, source=0, dest, tag=1, i;
typedef struct {
    float x, y, z;
    float velocity;
    int n, type;
} Particle;
Particle p[NELEM], particles[NELEM];
MPI_Datatype particletype, oldtypes[2];
int blockcounts[2];
MPI_Aint offsets[2];
```

Declaring the structure of the heterogeneous datatype
Float, Float, Float, Float, Int, Int

```
/* MPI_Aint type used to be consistent with syntax of */
/* MPI_Type_get_extnt routine */
MPI_Aint lb, float_extent;
```

```
MPI_Status stat;
```

```
MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
```

```
MPI_Type_get_extnt(MPI_FLOAT, &lb, &float_extent);
```

```
blockcounts[0] = 4; blockcounts[1] = 2;
offsets[0] = 0; offsets[1] = 4 * float_extent;
oldtypes[0] = MPI_FLOAT; oldtypes[1] = MPI_INT;
MPI_Type_create_struct(2, blockcounts, offsets, oldtypes,
&particletype);
MPI_Type_commit(&particletype);
```

```
if (rank == 0) {
    for (i=0; i<NELEM; i++) {
        particles[i].x = i * 1.0;
        particles[i].y = i * -1.0;
        particles[i].z = i * 1.0;
        particles[i].velocity = 0;
        particles[i].n = i;
        particles[i].type = i % 2;
```

Construct the heterogeneous datatype as an MPI datatype using Struct

```
    }
    for (i=0; i<numtasks; i++)
        MPI_Send(particles, NELEM, particletype, i, 0, MPI_COMM_WORLD);
    MPI_Recv(p, NELEM, particletype, 0, 0, MPI_COMM_WORLD, &stat);
    printf("rank= %d %3.2f %3.2f %3.2f\n", rank, p[3].y, p[3].z, p[3].velocity, p[3].n);
}
```

Populate the heterogeneous MPI datatype with heterogeneous data

```
MPI_Type_free(&particletype);
MPI_Finalize();
}
```


Example : Derived Datatype - struct

MPI_Type_struct

```
typedef struct { float x,y,z,velocity; int n,type; } Particle;  
Particle particles[NELEM];
```

```
MPI_Type_extent(MPI_FLOAT, &extent);
```

```
count = 2; oldtypes[0] = MPI_FLOAT; oldtypes[1] = MPI_INT;  
offsets[0] = 0; offsets[1] = 4 * extent;  
blockcounts[0] = 4; blockcounts[1] = 2;
```



particles[NELEM]

```
MPI_Type_struct(count, blockcounts, offsets, oldtypes, &particletype);
```

```
MPI_Send(particles, NELEM, particletype, dest, tag, comm);
```


Sends entire (NELEM) array of particles, each particle being comprised four floats and two integers.

```
[lsu00@master derived]$ qsub struct.pbs  
564.master.arete.cct.lsu.edu  
[lsu00@master derived]$ cat out.txt  
rank= 0    3.00 -3.00 3.00 0.25 3 1  
rank= 1    3.00 -3.00 3.00 0.25 3 1  
rank= 2    3.00 -3.00 3.00 0.25 3 1  
rank= 3    3.00 -3.00 3.00 0.25 3 1
```

Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- **Matrix-Vector multiplication : A Case Study**
- MPI Profiling calls
- Additional Topics
- Summary

Matrix Vector Multiplication



Slides for Parallel Programming Techniques & Applications Using Networked Workstations & Parallel Computers 2nd ed., by B. Wilkinson & M. Allen,
© 2004 Pearson Education Inc. All rights reserved.

Matrix Vector Multiplication


Multiplication of a matrix, **A** and a vector **B**, produces the vector **C** whose elements, c_i ($0 \leq i < n$), are computed as follows:

$$C_i = \sum_{k=0}^{k=m} A_{ik} * B_k$$

where **A** is an $n \times m$ matrix and **B** is a vector of size m and **C** is a vector of size n .

Matrix-Vector Multiplication

$$\mathbf{c} = \mathbf{A} \times \mathbf{b}$$



Example: Matrix-Vector Multiplication

```
#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
```

```
#define NRA 4      /* number of rows in matrix A */
#define NCA 4      /* number of columns in matrix A */
#define NCB 1      /* number of columns in matrix B */
#define MASTER 0   /* taskid of first task */
#define FROM_MASTER 1 /* setting a message type */
#define FROM_WORKER 2 /* setting a message type */
```

Define the dimensions of
the Matrix `a([4][4])` and
Vector `b([4][1])`

```
int main (int argc, char *argv[])
{
    int      numtasks,      /* number of tasks in partition */
            taskid,        /* a task identifier */
            numworkers,    /* number of worker tasks */
            source,        /* task id of message source */
            dest,          /* task id of message destination */
            mtype,         /* message type */
            rows,          /* rows of matrix A sent to each worker */
            averow, extra, offset, /* used to determine rows sent to each worker */
            i, j, k, rc;    /* misc */
}
```

Example: Matrix-Vector Multiplication

```
double    a[NRA][NCA],    /* Matrix A to be multiplied */
          b[NCA][NCB],    /* Vector B to be multiplied */
          c[NRA][NCB];    /* result Vector C */
```

```
MPI_Status status;
```

```
MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD,&taskid);
MPI_Comm_size(MPI_COMM_WORLD,&numtasks);
if (numtasks < 2 ) {
    printf("Need at least two MPI tasks. Quitting...\n");
    MPI_Abort(MPI_COMM_WORLD, rc);
    exit(1);
}
```

```
numworkers = numtasks-1;
```

```
/****** master task *****/
```

```
if (taskid == MASTER)
```

```
{
```

```
    printf("mpi_mm has started with %d tasks.\n",numtasks);
```

```
    printf("Initializing arrays...\n");
```

```
    for (i=0; i<NRA; i++)
```

```
        for (j=0; j<NCA; j++)
```

```
            a[i][j]= i+j;
```

```
    for (i=0; i<NCA; i++)
```

```
        for (j=0; j<NCB; j++)
```

```
            b[i][j]= (i+1)*(j+1);
```

Declare the matrix , vector
to be multiplied and the
resultant vector

MASTER Initializes the Matrix A :

0.00	1.00	2.00	3.00
1.00	2.00	3.00	4.00
2.00	3.00	4.00	5.00
3.00	4.00	5.00	6.00

MASTER Initializes B :

1.00
2.00
3.00
4.00

Example: Matrix-Vector Multiplication

```
for (i=0; i<NRA; i++)
{
    printf("\n");
    for (j=0; j<NCA; j++)
        printf("%6.2f ", a[i][j]);
}
for (i=0; i<NRA; i++)
{
    printf("\n");
    for (j=0; j<NCB; j++)
        printf("%6.2f ", b[i][j]);
}
```

Load Balancing : Dividing
the Matrix A based on the
number of processors

MASTER sends Matrix A to workers :

PROC[0] ::	0.00	1.00	2.00	3.00
PROC[1] ::	1.00	2.00	3.00	4.00
PROC[2] ::	2.00	3.00	4.00	5.00
PROC[3] ::	3.00	4.00	5.00	6.00

```
/* Send matrix data to the worker tasks */
```

```
averow = NRA/numworkers;
```

```
extra = NRA%numworkers;
```

```
offset = 0;
```

```
mtype = FROM_MASTER;
```

```
for (dest=1; dest<=numworkers; dest++)
```

```
{
```

```
    rows = (dest <= extra) ? averow+1 : averow;
```

```
printf("Sending %d rows to task %d offset=%d\n",rows,dest,offset);
```

```
MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
```

```
MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);
```

```
MPI_Send(&a[offset][0], rows*NCA, MPI_DOUBLE, dest, mtype,  
        MPI_COMM_WORLD);
```

```
MPI_Send(&b, NCA*NCB, MPI_DOUBLE, dest, mtype, MPI_COMM_WORLD);
```

```
offset = offset + rows;
```

```
}
```


MASTER Sends Vector B to Workers:

PROC[0] ::	1.00	2.00	3.00	4.00
PROC[1] ::	1.00	2.00	3.00	4.00
PROC[2] ::	1.00	2.00	3.00	4.00
PROC[3] ::	1.00	2.00	3.00	4.00

Example: Matrix-Vector Multiplication

```
/* Receive results from worker tasks */
mtype = FROM_WORKER;
for (i=1; i<=numworkers; i++)
{
    source = i;
    MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD, &status);
    MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD, &status);
    MPI_Recv(&c[offset][0], rows*NCB, MPI_DOUBLE, source, mtype,
             MPI_COMM_WORLD, &status);
    printf("Received results from task %d\n",source);
}

/* Print results */
printf("*****\n");
printf("Result Matrix:\n");
for (i=0; i<NRA; i++)
{
    printf("\n");
    for (j=0; j<NCB; j++)
        printf("%6.2f  ", c[i][j]);
}
printf("\n*****\n");
printf("Done.\n");
}
```



The Master process gathers the results and populates the result matrix in the correct order (easily done in this case because matrix index *i* is used to indicate position in result array)

Example: Matrix-Vector Multiplication

```
/****** worker task *****/
```

```
if (taskid > MASTER)
```

```
{
```

```
    mtype = FROM_MASTER;
```

```
    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
```

```
    MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD, &status);
```

```
    MPI_Recv(&a, rows*NCA, MPI_DOUBLE, MASTER, mtype, MPI_COMM_WORLD, &status);
```

```
    MPI_Recv(&b, NCA*NCB, MPI_DOUBLE, MASTER, mtype, MPI_COMM_WORLD, &status);
```

```
    for (k=0; k<NCB; k++)
```

```
        for (i=0; i<rows; i++)
```

```
        {
```

```
            c[i][k] = 0.0;
```

```
            for (j=0; j<NCA; j++)
```

```
                c[i][k] = c[i][k] + a[i][j] * b[j][k];
```

```
        }
```

```
    mtype = FROM_WORKER;
```

```
    MPI_Send(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
```

```
    MPI_Send(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD);
```

```
    MPI_Send(&c, rows*NCB, MPI_DOUBLE, MASTER, mtype, MPI_COMM_WORLD);
```

```
}
```

```
MPI_Finalize();
```

```
}
```

Worker Processes receive workload

Proc[1] A : 1.00 2.00 3.00 4.00

Proc[1] B : 1.00 2.00 3.00 4.00

Calculate Result

Proc[1] C: 1.00 + 4.00 + 9.00 + 16.00

Worker sends result to MASTER

Proc[1] C : 20.00

Example: Matrix-Vector Multiplication (Results)

```
[lsu00@master mm]$ qsub mm.pbs
568.master.arete.cct.lsu.edu
[lsu00@master mm]$ cat out.txt
mpi_mm has started with 4 tasks.
Initializing arrays...
```

```
*****
```

0.00	1.00	2.00	3.00
1.00	2.00	3.00	4.00
2.00	3.00	4.00	5.00
3.00	4.00	5.00	6.00

```
*****
```

1.00
2.00
3.00
4.00

```
*****
```

```
Sending 2 rows to task 1 offset=0
Sending 1 rows to task 2 offset=2
Sending 1 rows to task 3 offset=3
Received results from task 1
Received results from task 2
Received results from task 3
```

```
*****
```

Result Matrix:

20.00
30.00
40.00
50.00

```
*****
```

Done.

Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- **MPI Profiling calls**
- Additional Topics
- Summary

MPI Profiling : MPI_Wtime

Function: `MPI_Wtime()`

`double MPI_Wtime()`

Description:

Returns time in seconds elapsed on the calling processor. Resolution of time scale is determined by the MPI environment variable `MPI_WTICK`. When the MPI environment variable `MPI_WTIME_IS_GLOBAL` is defined and set to true, the the value of `MPI_Wtime` is synchronized across all processes in `MPI_COMM_WORLD`

```
double time0;  
...  
time0 = MPI_Wtime();  
...  
printf("Hello From Worker #%d %lf \n", rank, (MPI_Wtime() - time0));
```

http://www-unix.mcs.anl.gov/mpi/www/www3/MPI_Wtime.html

Timing Example: MPI_Wtime

```
#include <stdio.h>
#include "mpi.h"

main(int argc, char **argv)
{
    int size, rank;
    double time0, time1;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    time0 = MPI_Wtime();

    if(rank==0)
    {
        printf(" Hello From Proc0 Time = %lf \n", (MPI_Wtime() - time0));
    }
    else
    {
        printf("Hello From Worker #%d %lf \n", rank, (MPI_Wtime() - time0));
    }
    MPI_Finalize();
}
```

Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- **Additional Topics**
- Summary

Potential Pitfall

- We need to be sure that every receive has a matching send
 - Suppose process q calls
 - `MPI_Send(send_buf_p, send_buf_sz, send_type, dest, send_tag, send_comm);`
 - And process r calls
 - `MPI_Recv(recv_buf_p, recv_buf_sz, recv_type, src, recv_tag, recv_comm, &status);`
 - The message sent by q with the above call to `MPI_Send` can be received by r with the call to `MPI_Recv` if:
 - `recv_comm = send_comm,`
 - `recv_tag = send_tag,`
 - `dest = r,` and
 - `src = q`

Wildcard Arguments

- MPI_ANY_SOURCE
- MPI_ANY_TAG
- Only a receiver can use a wildcard argument. Senders must specify a process rank and a nonnegative tag
- Example:
 - If process 0 executes the following code, it can receive the results in the order in which the processes finish:

```
for (i = 1; i < comm_sz; i++) {  
    MPI_Recv(result, result_sz, result_type, MPI_ANY_SOURCE,  
             result_tag, comm, MPI_STATUS_IGNORE);  
    Process_result(result);  
}
```

Input/Output

- Although the MPI standard doesn't specify which processes have access to which I/O devices, virtually all MPI implementations allow *all* the processes in MPI_COMM_WORLD full access to stdout and stderr
 - so most MPI implementations allow all processes to execute `printf` and `fprintf(stderr, ...)`
- Unlike output, most MPI implementations only allow process 0 in MPI_COMM_WORLD access to stdin
 - This makes sense: If multiple processes have access to stdin, which process should get which parts of the input data?

Collective vs. Point-to-point Communications

- All the processes in the communicator must call the same collective function
 - For example, a program that attempts to match a call to `MPI_Reduce` on one process with a call to `MPI_Recv` on another process is erroneous, and, in all likelihood, the program will hang or crash.
- The arguments passed by each process to an MPI collective communication must be “compatible.”
 - For example, if one process passes in 0 as the dest process and another passes in 1, then the outcome of a call to `MPI_Reduce` is erroneous, and, once again, the program is likely to hang or crash.
- Collective communications don't use tags

Collective vs. Point-to-point Communications (Cont'd)

- The names of the memory locations are irrelevant to the matching, of the calls to MPI_Reduce
 - The order of the calls will determine the matching
 - For example, Suppose that each process calls MPI_Reduce with operator MPI_SUM, and destination process 0
 - The value stored in b will be $1 + 2 + 1 = 4$
 - The value stored in d will be $2 + 1 + 2 = 5$

Table 3.3 Multiple Calls to MPI_Reduce

Time	Process 0	Process 1	Process 2
0	a = 1; c = 2	a = 1; c = 2	a = 1; c = 2
1	MPI_Reduce(&a, &b, ...)	MPI_Reduce(&c, &d, ...)	MPI_Reduce(&a, &b, ...)
2	MPI_Reduce(&c, &d, ...)	MPI_Reduce(&a, &b, ...)	MPI_Reduce(&c, &d, ...)

MPI topics we did Not Cover

- Topologies: map a communicator onto, say, a 3D Cartesian processor grid
 - Implementation can provide ideal logical to physical mapping
- Rich set of I/O functions: individual, collective, blocking and non-blocking
 - Collective I/O can lead to many small requests being merged for more efficient I/O
- One-sided communication: puts and gets with various synchronization schemes
 - Implementations not well-optimized and rarely used
 - Redesign of interface is underway
- Task creation and destruction: change number of tasks during a run
 - Few implementations available

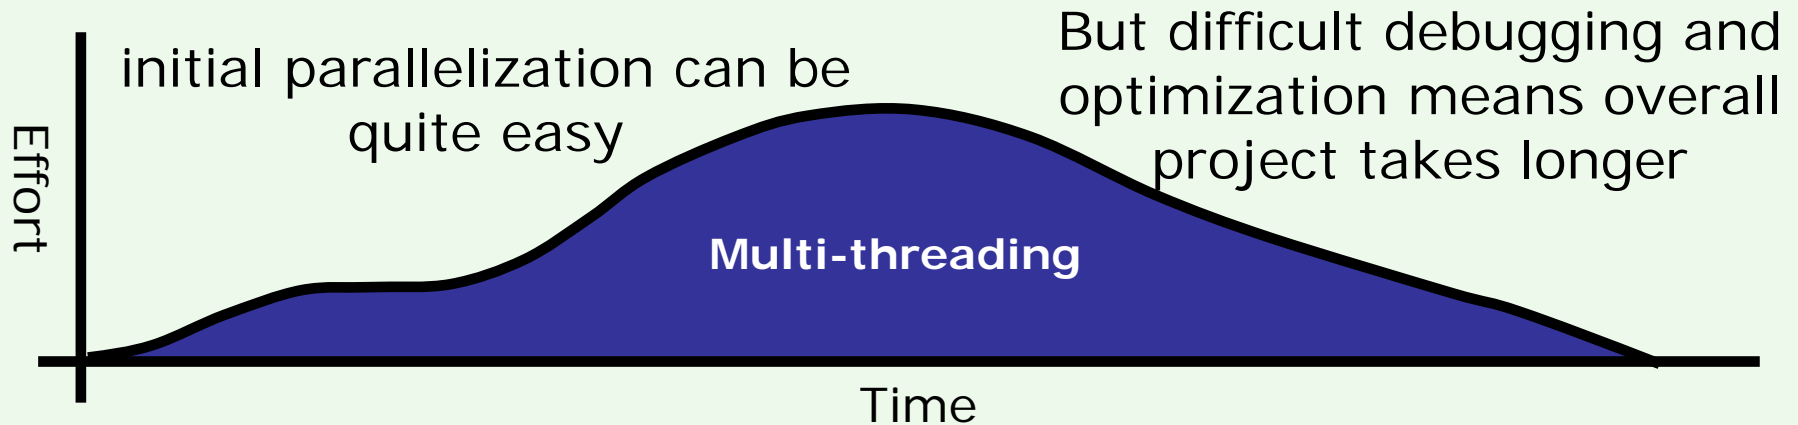
Topics

- MPI Collective Calls: Synchronization Primitives
- MPI Collective Calls: Communication Primitives
- MPI Collective Calls: Reduction Primitives
- Derived Datatypes: Introduction
- Derived Datatypes: Contiguous
- Derived Datatypes: Vector
- Derived Datatypes: Indexed
- Derived Datatypes: Struct
- Matrix-Vector multiplication : A Case Study
- MPI Profiling calls
- Additional Topics
- Summary

MPI isn't as hard as many believe

- There are over 330 functions in the MPI spec, but most programs only use a small subset:
 - Startup
 - MPI_Init, MPI_Finalize
 - Information on the processes
 - MPI_Comm_rank, MPI_Comm_size,
 - Point-to-point communication
 - MPI_Irecv, MPI_Isend, MPI_Wait, MPI_Send, MPI_Recv
 - Collective communication
 - MPI_Reduce, MPI_Allreduce, MPI_Bcast, MPI_Allgather

Multithreading vs Message Passing



Proving that a shared address space program using semaphores is race free is an NP-complete problem

P. N. Klein, H. Lu, and R. H. B. Netzer, Detecting Race Conditions in Parallel Programs that Use Semaphores, Algorithmica, vol. 35 pp. 321-345, 2003

