
Parallel Programming

Hadoop and MapReduce Programming

Professor Yi-Ping You (游逸平)

Department of Computer Science

<http://www.cs.nctu.edu.tw/~ypyou/>



Acknowledgement

- This set of slides is adapted from lectures by
 - ✦ Prof. Jimmy Lin (University of Maryland), “Introduction to Cloud Computing”
 - ✦ Prof. Anand Rajaraman (Stanford University), “Data Mining”
 - ✦ Prof. Adriana Iamnitchi (University of South Florida), “Basics of Parallel and Distributed Systems”



Outline

- Introduction to MapReduce
 - ◆ Google File System
- The Hadoop Framework



What is MapReduce?

- Data-parallel programming model for clusters of commodity machines
- Pioneered by Google (published in OSDI'04)
 - ✦ Processes 20 PB of data per day
 - ◆ Index building for Google Search
 - ◆ Article clustering for Google News
 - ◆ Statistical machine translation
- Popularized by Apache Hadoop project
 - ✦ Used by Yahoo!, Facebook, Amazon, ...
 - ◆ Index building for Yahoo! Search
 - ◆ Spam detection for Yahoo! Mail
 - ◆ Ad optimization for Facebook




Commodity Clusters

- Web data sets can be very large
 - Tens to hundreds of terabytes
- Cannot mine on a single server (why?)
- Standard architecture emerging:
 - Cluster of commodity Linux nodes
 - Gigabit ethernet interconnect
- How to organize computations on this architecture?
 - Mask issues such as hardware failure

Cluster Architecture

2-10 Gbps backbone between racks

1 Gbps between
any pair of nodes
in a rack



Each rack contains 16-64 nodes

Stable storage

- First order problem: if nodes can fail, how can we store data persistently?
- Answer: Distributed File System
 - Provides global file namespace
 - Google GFS; Hadoop HDFS; Kosmix KFS
- Typical usage pattern
 - Huge files (100s of GB to TB)
 - Data is rarely updated in place
 - Reads and appends are common

Distributed File System

□ Chunk Servers

- a.k.a. Data Nodes in HDFS
- File is split into contiguous chunks
- Typically each chunk is 16-64MB
- Each chunk replicated (usually 2x or 3x)
- Try to keep replicas in different racks

□ Master node

- a.k.a. Name Node in HDFS
- Stores metadata
- Might be replicated

□ Client library for file access

- Talks to master to find chunk servers
- Connects directly to chunkservers to access data

Motivation for MapReduce

- Large-Scale Data Processing
 - Want to use 1000s of CPUs
 - But don't want hassle of *managing* things

- MapReduce Architecture provides
 - Automatic parallelization & distribution
 - Fault tolerance
 - I/O scheduling
 - Monitoring & status updates

MapReduce Goals

□ Cloud Environment:

- Commodity nodes (cheap, but unreliable)
- Commodity network (low bandwidth)
- Automatic fault-tolerance (fewer admins)

□ Scalability to large data volumes:

- Scan 100 TB on 1 node @ 50 MB/s = 24 days
- Scan on 1000-node cluster = 35 minutes

What is Map/Reduce

- Map/Reduce
 - Programming model from LISP
 - (and other functional languages)
- Many problems can be phrased this way
- Easy to distribute across nodes
- Nice retry/failure semantics

Map in LISP (Scheme)

□ (map *f list* [*list*₂ *list*₃ ...])


□ (map square '(1 2 3 4))
■ (1 4 9 16)

Unary operator



Map

- Map is a higher-order function
- How map works:
 - Function is applied to every element in a list
 - Result is a new list
- No limit to map parallelization since maps are independent



Reduce in LISP (Scheme)

□ (reduce ***f id list***)

Binary operator

□ (reduce + 0 '(1 4 9 16))

■ (+ 16 (+ 9 (+ 4 (+ 1 0))))


■ 30

□ (reduce + 0
 (map square (l₁ l₂)))

Reduce

- Reduce is also a higher-order function
- How reduce works:
 - Accumulator set to initial value
 - Function applied to list element and the accumulator
 - Result stored in the accumulator
 - Repeated for every item in the list
 - Result is the final value in the accumulator

Reduce



- We can reorder folding if the reduce function is commutative and associative

Typical Problem

- ☐ Iterate over a large number of records
- ☐ **Map** Extract something of interest from each
- ☐ Shuffle and sort intermediate results
- ☐ Aggregate intermediate results **Reduce**
- ☐ Generate final output

Key idea: provide an abstraction at the point of these two operations

MapReduce

- Programmers specify two functions:

Map function:

$$(K_{in}, V_{in}) \rightarrow \text{list}\langle(K_{inter}, V_{inter})\rangle$$

Reduce function:

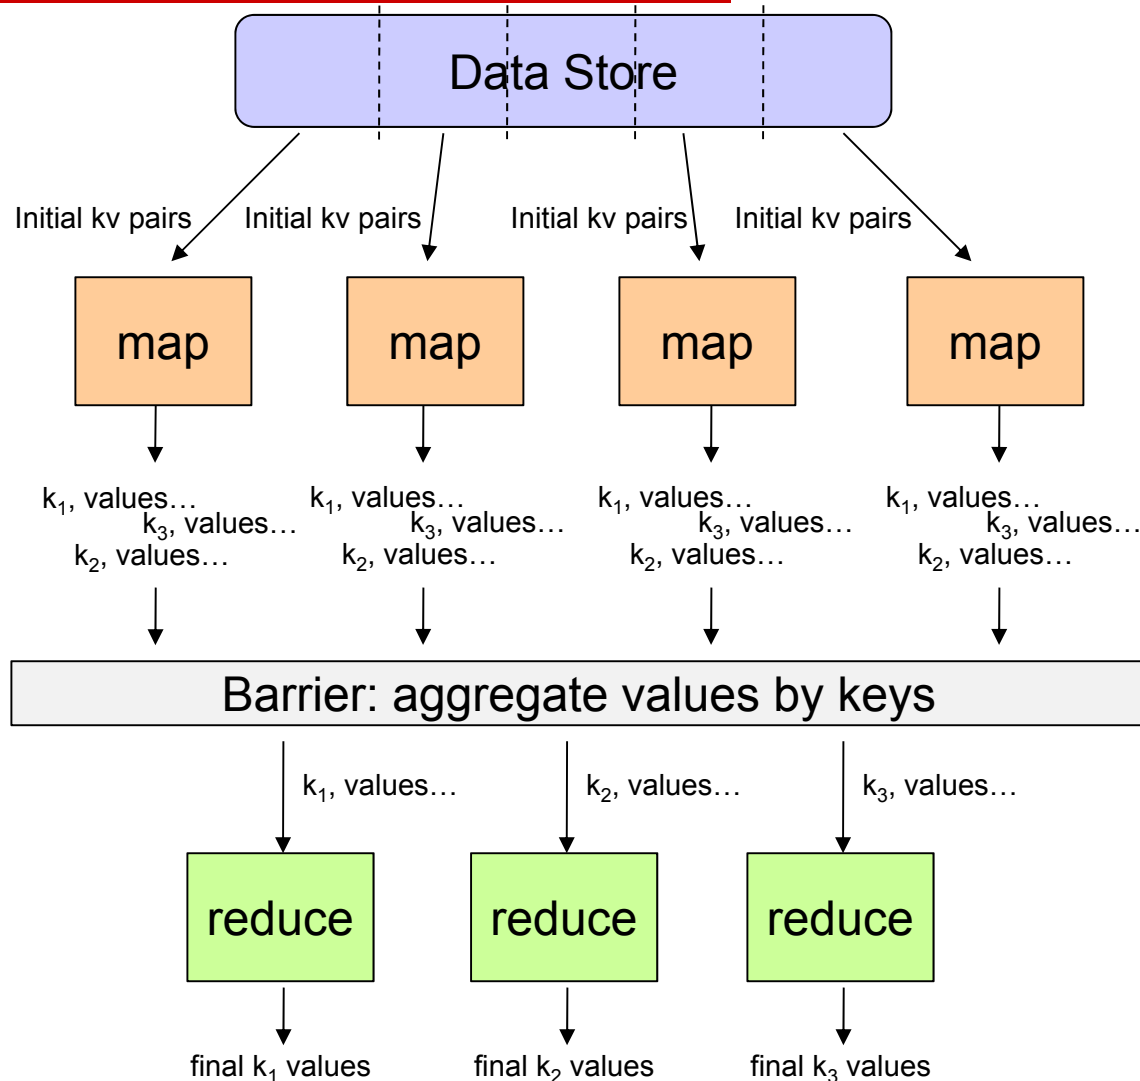
$$(K_{inter}, \text{list}\langle V_{inter} \rangle) \rightarrow \text{list}\langle(K_{out}, V_{out})\rangle$$

- Usually, programmers also specify:

partition (k' , number of partitions) \rightarrow partition for k'

- Often a simple hash of the key, e.g. $\text{hash}(k') \bmod n$
- Allows reduce operations for different keys in parallel

It's just divide and conquer



“Hello World”: Word Count

- We have a large file of words
- Count the number of times each distinct word appears in the file
- *Sample application:* analyze web server logs to find popular URLs

Word Count (2)

- ❑ Case 1: Entire file fits in memory
- ❑ Case 2: File too large for mem, but all `<word, count>` pairs fit in mem
- ❑ Case 3: File on disk, too many distinct words to fit in memory
 - **`xargs -n1 < datafile | sort | uniq -c`**

Word Count (3)

- ❑ To make it slightly harder, suppose we have a large corpus of documents
- ❑ Count the number of times each distinct word occurs in the corpus

```
words(docs/*) | sort | uniq -c
```

where **words** takes a file and outputs the words in it, one to a line
- ❑ The above captures the essence of MapReduce
 - Great thing is that it is naturally parallelizable

MapReduce

- Input: a set of key/value pairs
- User supplies two functions:
 - $\text{map}(k, v) \rightarrow \text{list}(k1, v1)$
 - $\text{reduce}(k1, \text{list}(v1)) \rightarrow v2$
- $(k1, v1)$ is an intermediate key/value pair
- Output is the set of $(k1, v2)$ pairs

Word Count using MapReduce

```
map(key, value):
```

```
// key: document name; value: text of document
```

```
  for each word w in value:
```

```
    emit(w, 1)
```

```
reduce(key, values):
```

```
// key: a word; values: an iterator over counts
```

```
  result = 0
```

```
  for each count v in values:
```

```
    result += v
```

```
  emit(key, result)
```


Count, Illustrated

map(key=url, val=contents):


For each word w in contents, emit (w , "1")

reduce(key=word,
values=uniq_counts):


Sum all "1" s in values list

Emit result "(word, sum)"

see bob run
see spot throw

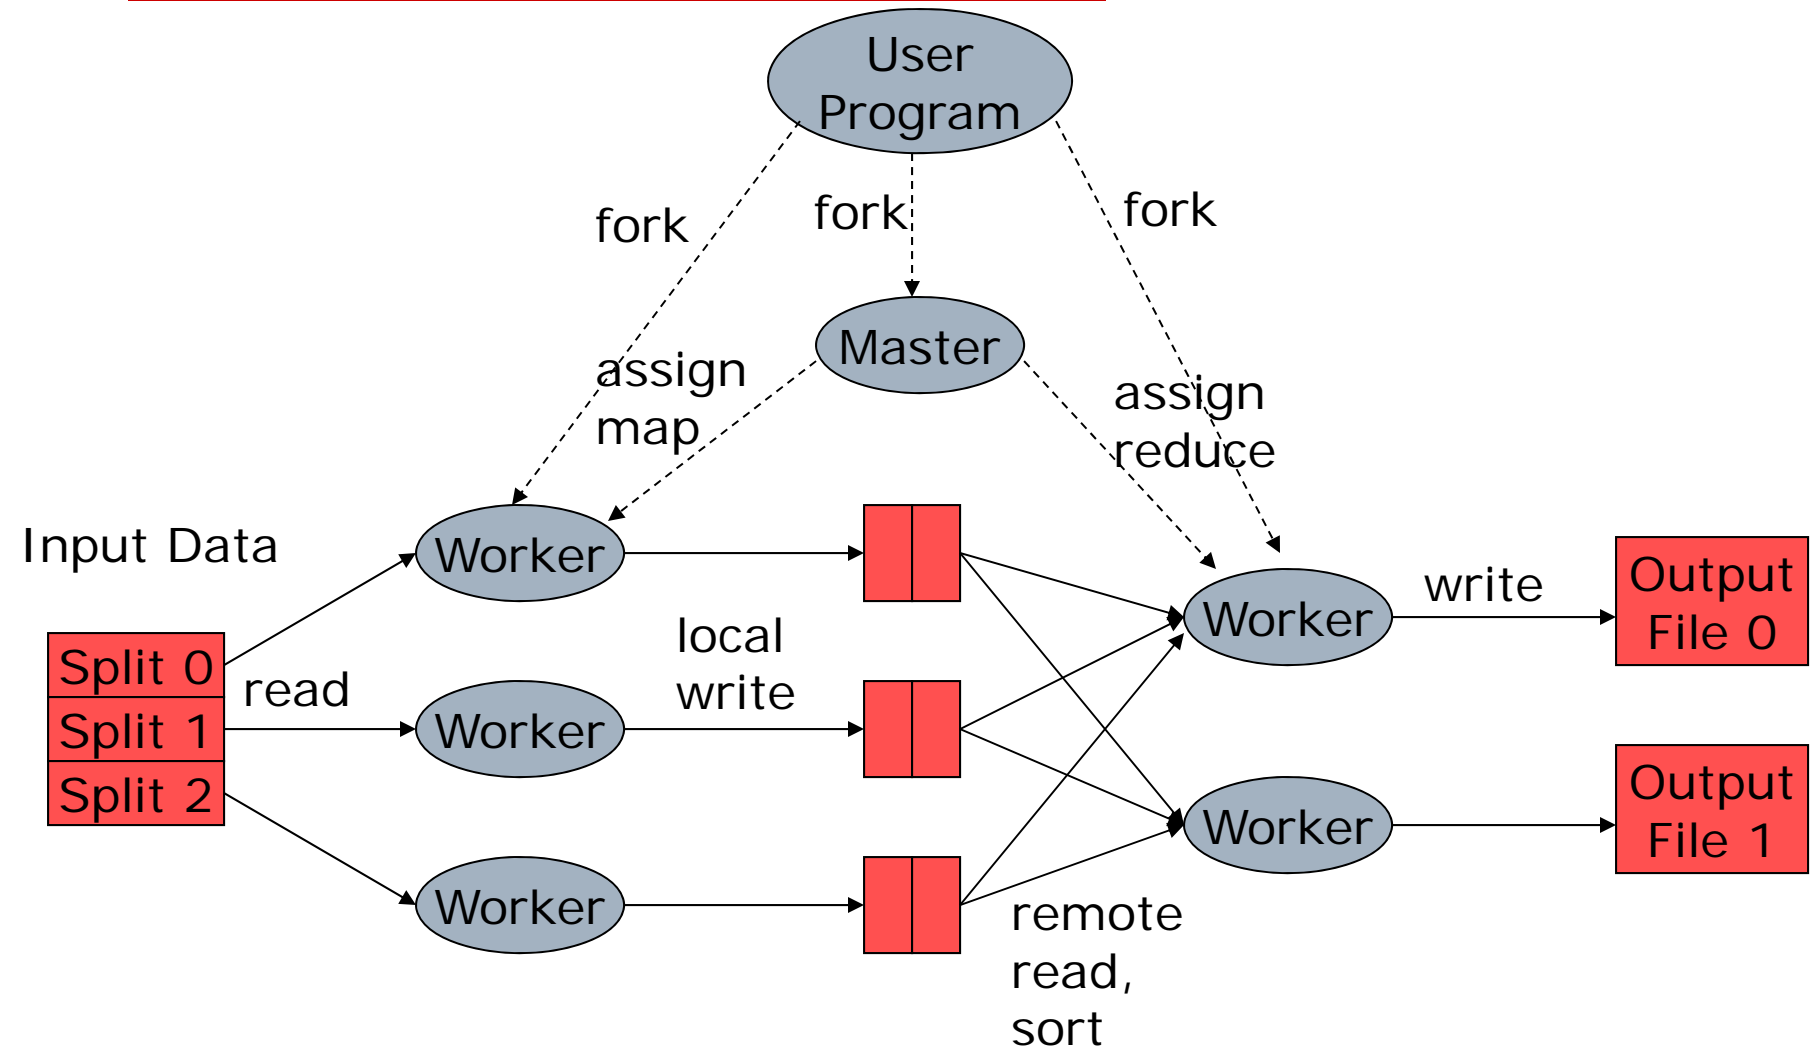


see	1
bob	1
run	1
see	1
spot	1
throw	1



bob	1
run	1
see	2
spot	1
throw	1

Behind the scenes...



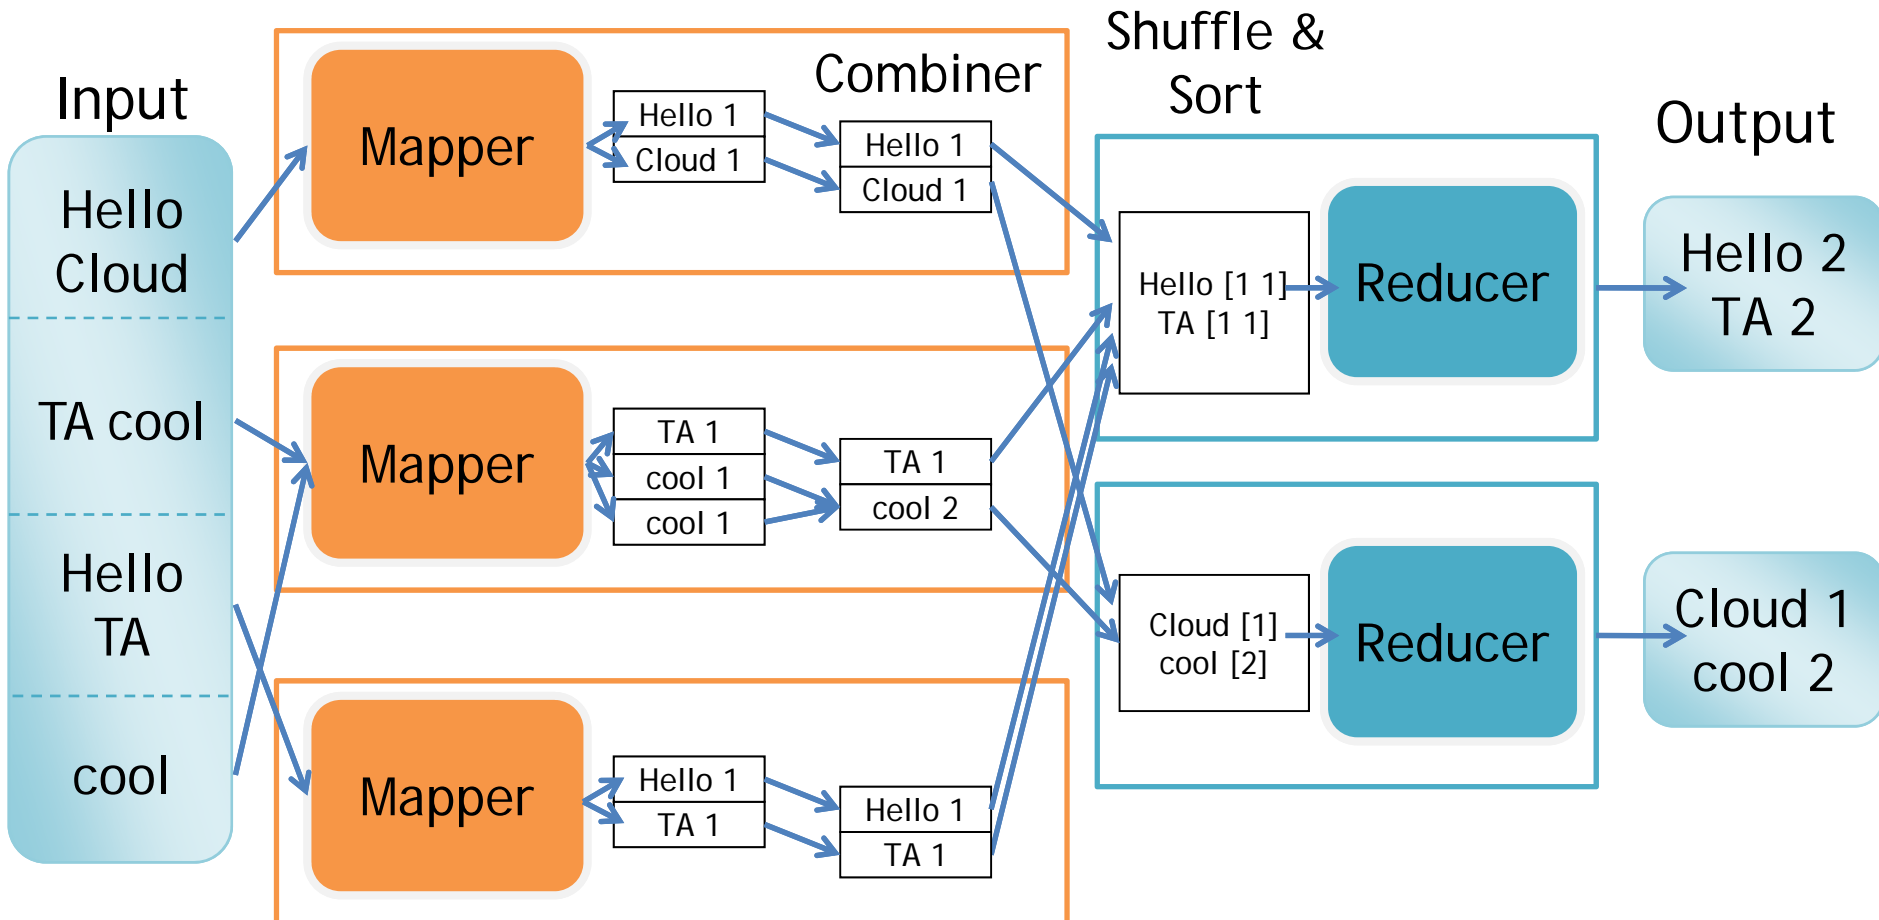
Bandwidth Optimization

- **Issue:** large number of key-value pairs
- **Solution:** use “Combiner” functions
 - Executed on same machine as mapper
 - Results in a “mini-reduce” right after the map phase
 - Reduces key-value pairs to save bandwidth

Example: Word Count

```
def map(line_num, line):  
    foreach word in line.split():  
        output(word, 1)
```

```
def reduce(word, counts):  
    output(word, sum(counts))
```



Combiners


- ❑ Often a map task will produce many pairs of the form $(k, v1), (k, v2), \dots$ for the same key k
 - E.g., popular words in Word Count
- ❑ Can save network time by pre-aggregating at mapper
 - `combine(k1, list(v1)) → v2`
 - Usually same as reduce function
- ❑ Works only if reduce function is commutative and associative

Skew Problem

- **Issue:** reduce is only as fast as the slowest map
- **Solution:** redundantly execute map operations, use results of first to finish
 - Addresses hardware problems...
 - But not issues related to inherent distribution of data

Model is Widely Applicable

MapReduce Programs In Google Source Tree



Example uses:

distributed grep

term-vector / host

document clustering

distributed sort

web access log stats

machine learning

web link-graph reversal

inverted index construction

statistical machine
translation

Example: Reverse Web-Link Graph

- ❑ Find all pages that link to a certain page
- ❑ Map Function
 - Outputs $\langle target, source \rangle$ pairs for each link to a *target* URL found in a *source* page
 - For each page we know what pages it links to
- ❑ Reduce Function
 - Concatenates the list of all *source* URLs associated with a given *target* URL and emits the pair: $\langle target, list(source) \rangle$
 - For a given web page, we know what pages link to it


What about the following problems?

- ☐ How do we assign work units to workers?
- ☐ What if we have more work units than workers?
- ☐ What if workers need to share partial results?
- ☐ How do we aggregate partial results?
- ☐ How do we know all the workers have finished?
- ☐ What if workers die?

MapReduce Runtime

- Handles scheduling
 - Assigns workers to map and reduce tasks
- Handles “data distribution”
 - Moves the process to the data
- Handles synchronization
 - Gathers, sorts, and shuffles intermediate data
- Handles faults
 - Detects worker failures and restarts
- Everything happens on top of a distributed file system

How do we get data to the workers?



What's the problem here?



Distributed File System

- Don't move data to workers... Move workers to the data!
 - ✦ Store data on the local disks for nodes in the cluster
 - ✦ Start up the workers on the node that has the data local
- Why?
 - ✦ Not enough RAM to hold all the data in memory
 - ✦ Disk access is slow, disk throughput is good
- A distributed file system is the answer
 - ✦ GFS (Google File System)
 - ✦ HDFS for Hadoop



Outline

- Introduction to MapReduce
 - ✦ Google File System
- The Hadoop Framework



GFS: Assumptions

- Commodity hardware over “exotic” hardware
- High component failure rates
 - ✦ Inexpensive commodity components fail all the time
- “Modest” number of HUGE files
- Files are write-once, mostly appended to
- Perhaps concurrently
- Large streaming reads over random access
- High sustained throughput over low latency

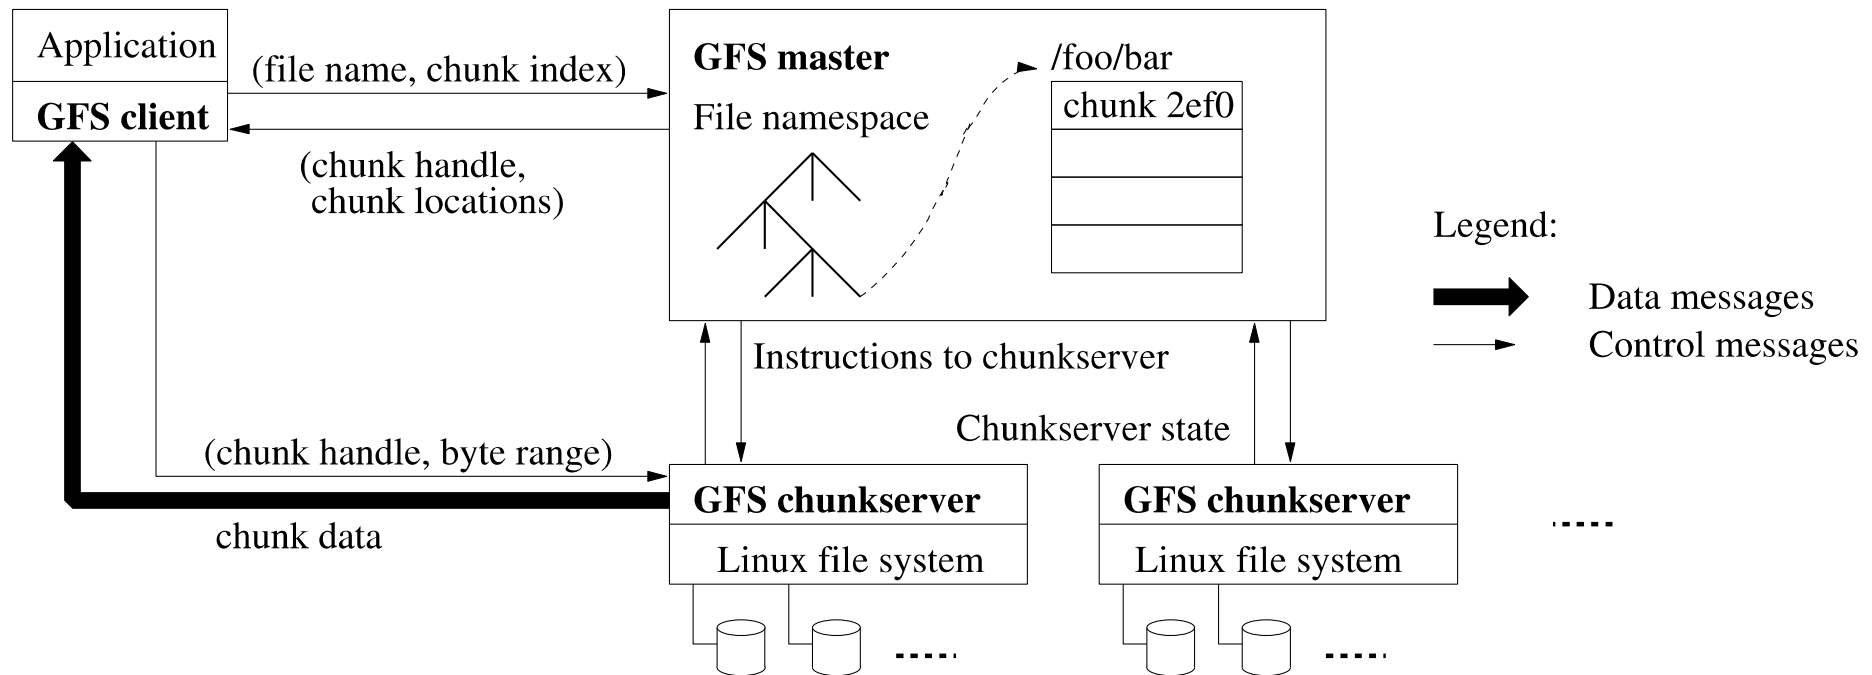


GFS: Design Decisions

- Files stored as chunks
 - ✦ Fixed size (64MB)
- Reliability through replication
 - ✦ Each chunk replicated across 3+ chunkservers
- Single master to coordinate access, keep metadata
 - ✦ Simple centralized management
- No data caching
 - ✦ Little benefit due to large data sets, streaming reads
- Simplify the API
 - ✦ Push some of the issues onto the client



GFS Architecture



Source: Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. 2003. The Google file system, SOSP 2003.



Single Master

- We know this is a:
 - ✦ Single point of failure
 - ✦ Scalability bottleneck
- GFS solutions:
 - ✦ Shadow masters
 - ✦ Minimize master involvement
 - ◆ Never move data through it, use only for metadata (and cache metadata at clients)
 - ◆ Large chunk size
 - ◆ Master delegates authority to primary replicas in data mutations (chunk leases)
- Simple, and good enough!



Master's Responsibilities (1/2)

- Metadata storage
- Namespace management/locking
- Periodic communication with chunkservers
 - ✦ Give instructions, collect state, track cluster health
- Chunk creation, re-replication, rebalancing
 - ✦ Balance space utilization and access speed
 - ✦ Spread replicas across racks to reduce correlated failures
 - ✦ Re-replicate data if redundancy falls below threshold
 - ✦ Rebalance data to smooth out storage and request load



Master's Responsibilities (2/2)

■ Garbage Collection

- ✦ Simpler, more reliable than traditional file delete
- ✦ Master logs the deletion, renames the file to a hidden name
- ✦ Lazily garbage collects hidden files

■ Stale replica deletion

- ✦ Detect “stale” replicas using chunk version numbers



Metadata

- Global metadata is stored on the master
 - ✦ File and chunk namespaces
 - ✦ Mapping from files to chunks
 - ✦ Locations of each chunk's replicas
- All in memory (64 Mbytes / chunk)
 - ✦ Fast
 - ✦ Easily accessible
- Master has an operation log for persistent logging of critical metadata updates
 - ✦ Persistent on local disk
 - ✦ Replicated
 - ✦ Checkpoints for faster recovery



Coordination

- Master data structures
 - Task status: (idle, in-progress, completed)
 - Idle tasks get scheduled as workers become available
 - When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
 - Master pushes this info to reducers
- Master pings workers periodically to detect failures

Failures

❑ Map worker failure

- Map tasks completed or in-progress at worker are reset to idle
- Reduce workers are notified when task is rescheduled on another worker

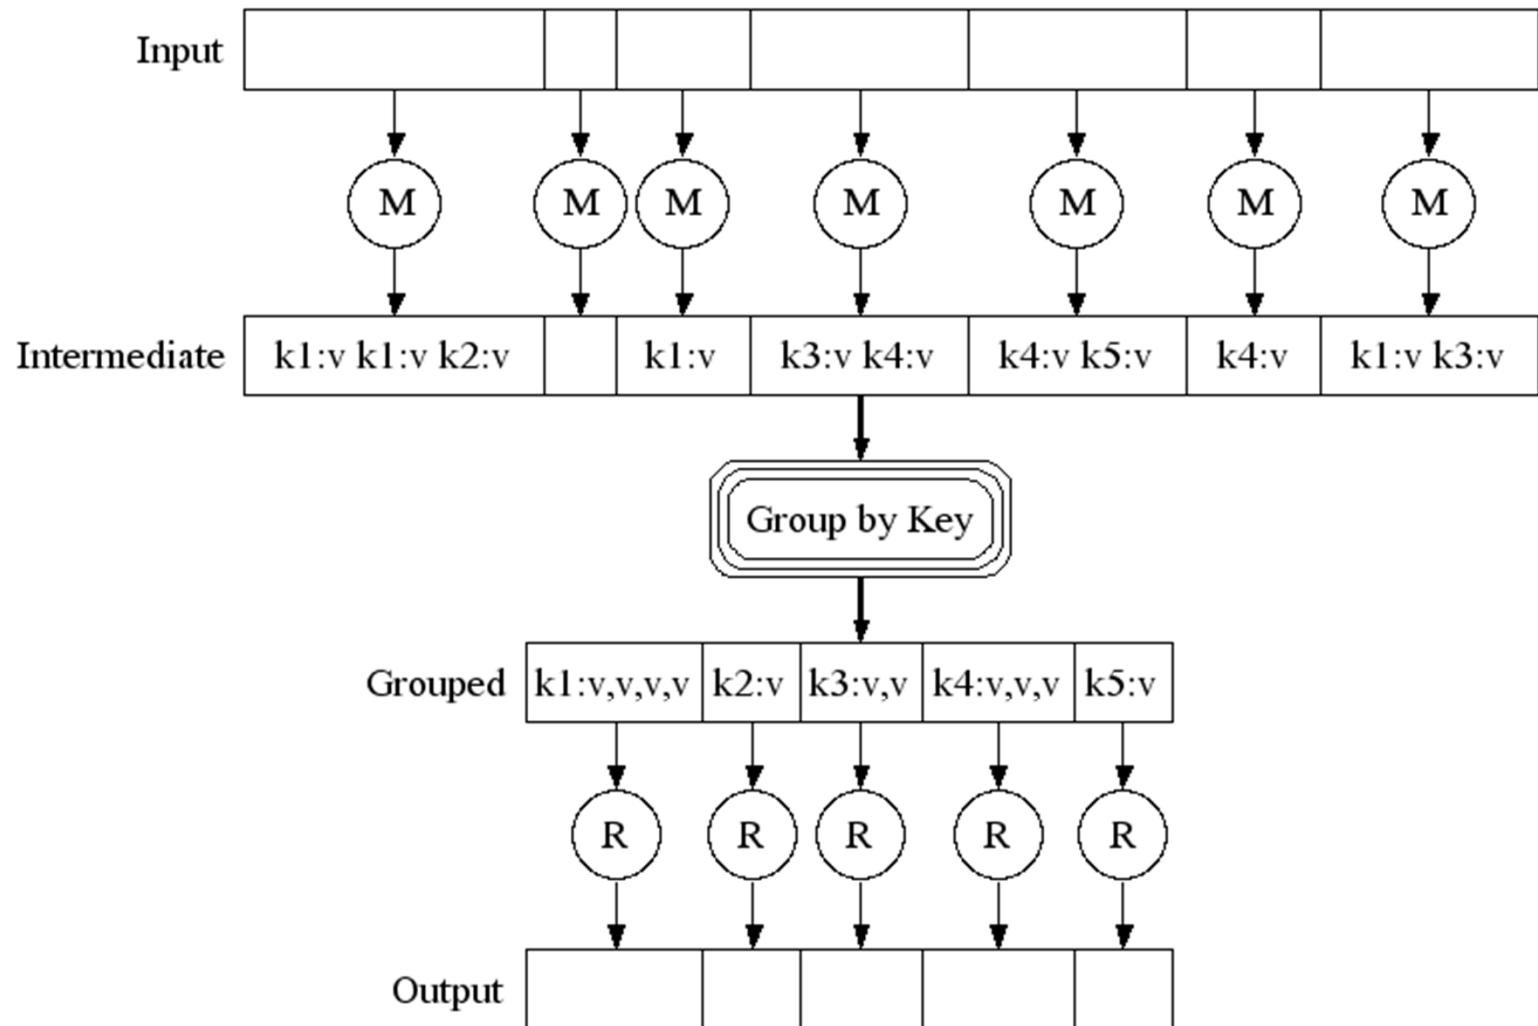
❑ Reduce worker failure

- Only in-progress tasks are reset to idle

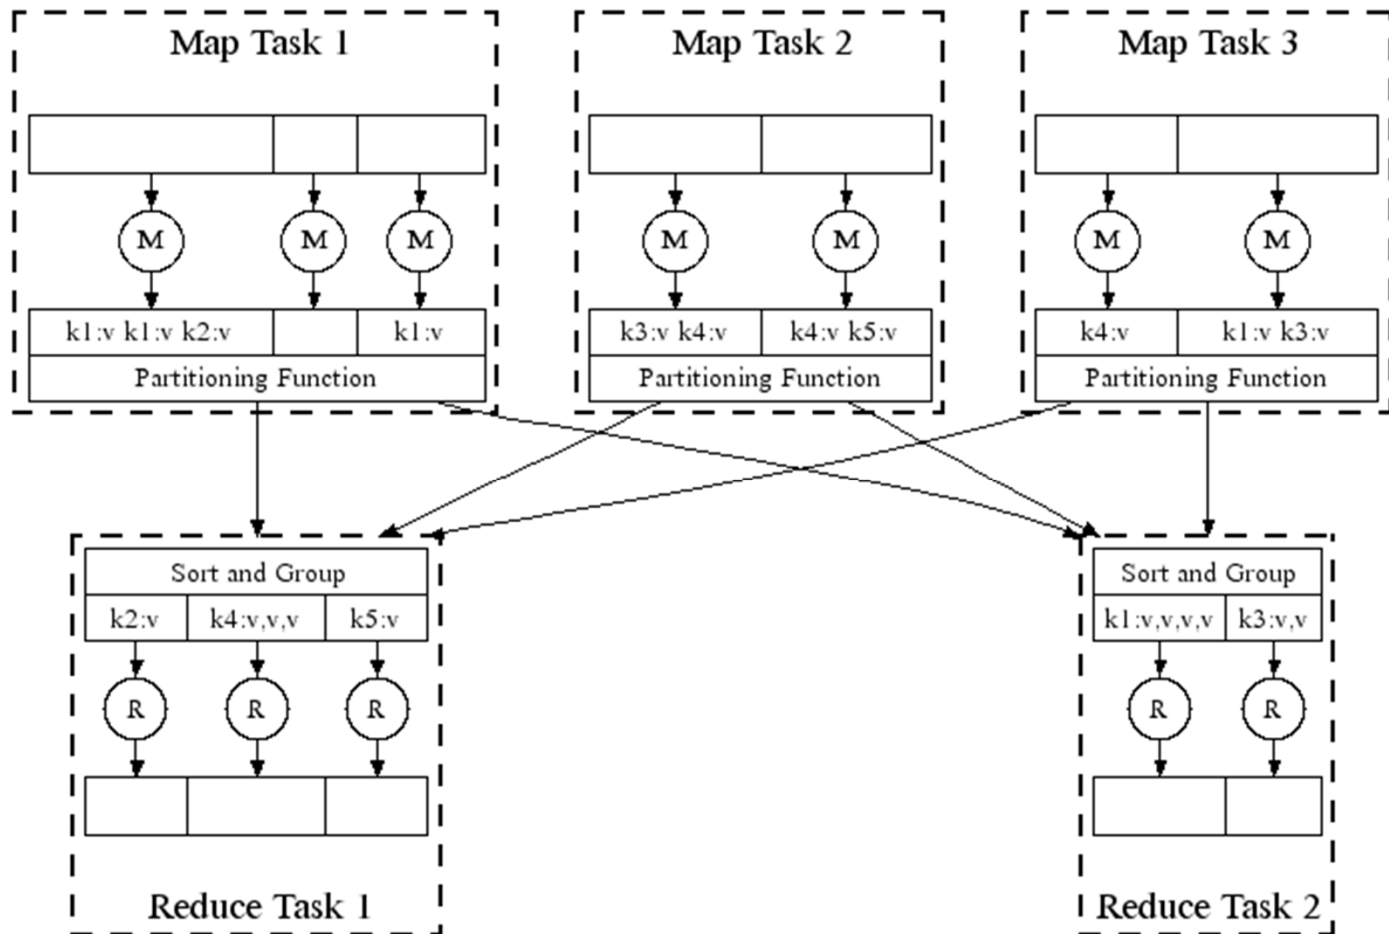
❑ Master failure

- MapReduce task is aborted and client is notified

Execution



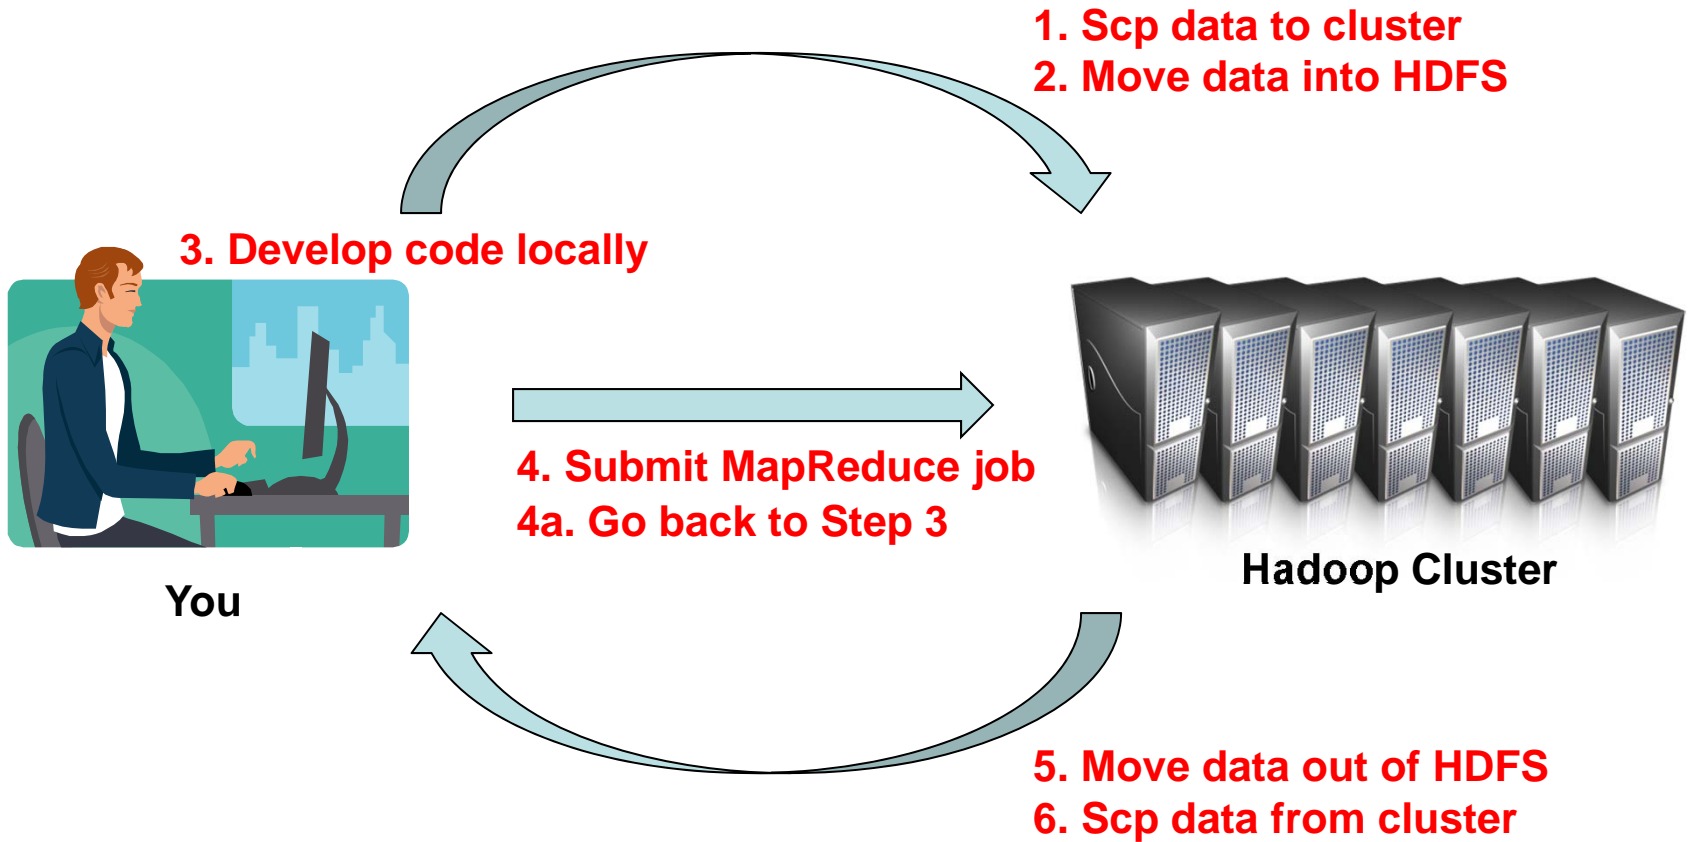
Parallel Execution



How many Map and Reduce jobs?

- M map tasks, R reduce tasks
- Rule of thumb:
 - Make M and R much larger than the number of nodes in cluster
 - One DFS chunk per map is common
 - Improves dynamic load balancing and speeds recovery from worker failure
- Usually R is smaller than M, because output is spread across R files

From Theory to Practice



Execution Summary

- ❑ How is this distributed?
 1. Partition input key/value pairs into chunks, run `map()` tasks in parallel
 2. After all `map()`s are complete, consolidate all emitted values for each unique emitted key
 3. Now partition space of output map keys, and run `reduce()` in parallel
- ❑ If `map()` or `reduce()` fails, reexecute!

Outline

- Introduction to MapReduce
 - ◆ Google File System
- The Hadoop Framework



Hadoop

- Hadoop (High-availability distributed object-oriented platform)
- An open-source software framework that supports data-intensive distributed applications
- Distributed File System (HDFS)
 - ✦ Runs on same machines
 - ✦ Replicates data 3x for fault-tolerance
- It supports the running of applications on large clusters of commodity hardware

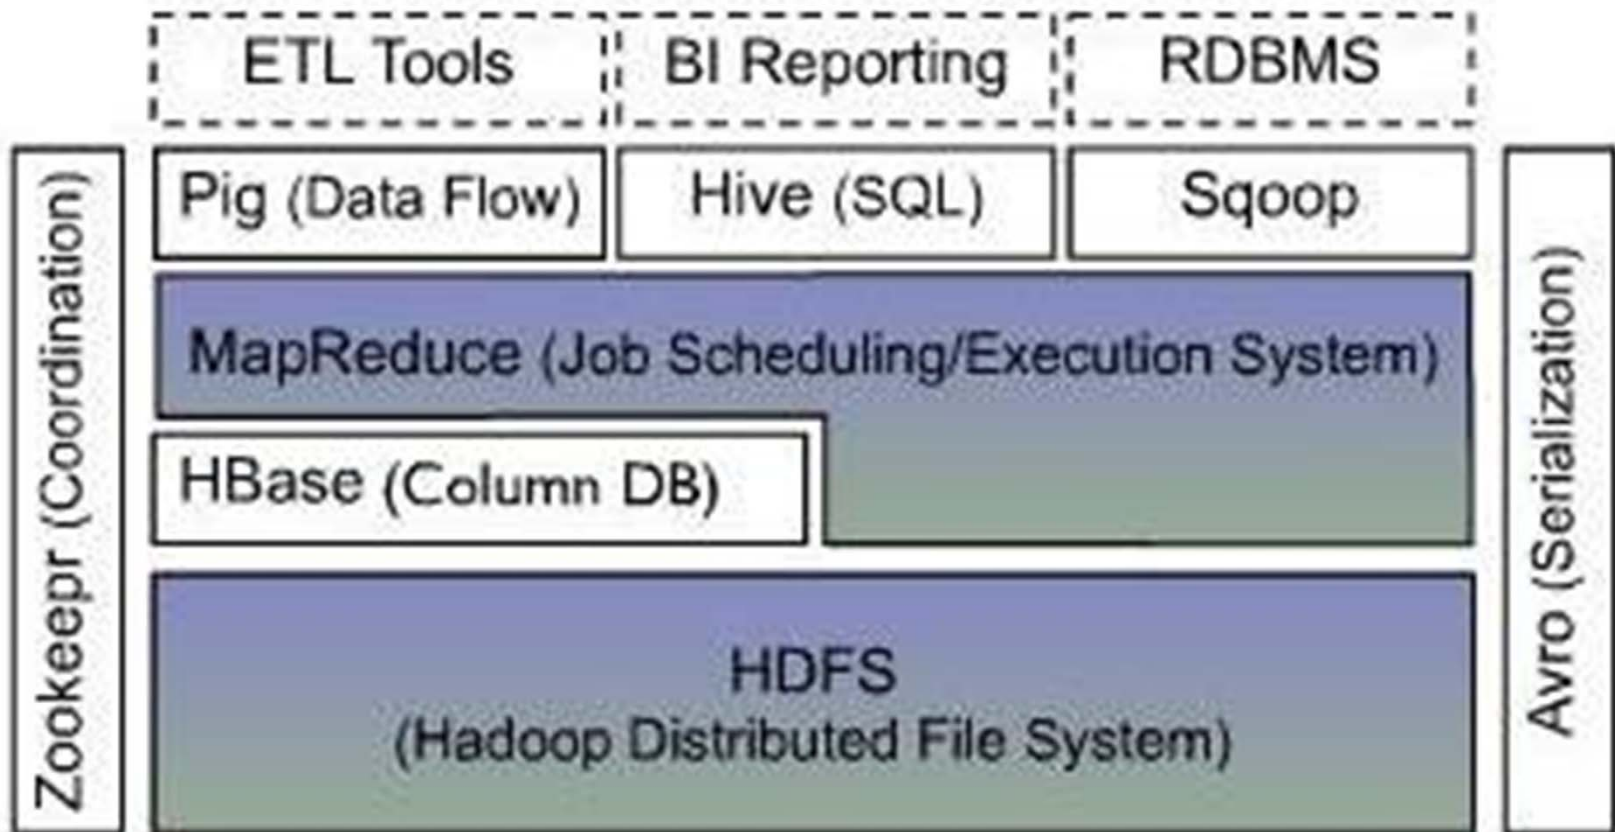


Hadoop

- Hadoop was derived from Google's MapReduce and Google File System (GFS) papers
- Hadoop is written in the Java programming language
 - ✦ Download from: <http://hadoop.apache.org/>



Hadoop Ecosystem



Source: Cloudera



Terminology

Google calls it:	Hadoop equivalent:
MapReduce	Hadoop
GFS	HDFS
Bigtable	HBase
Chubby	Zookeeper

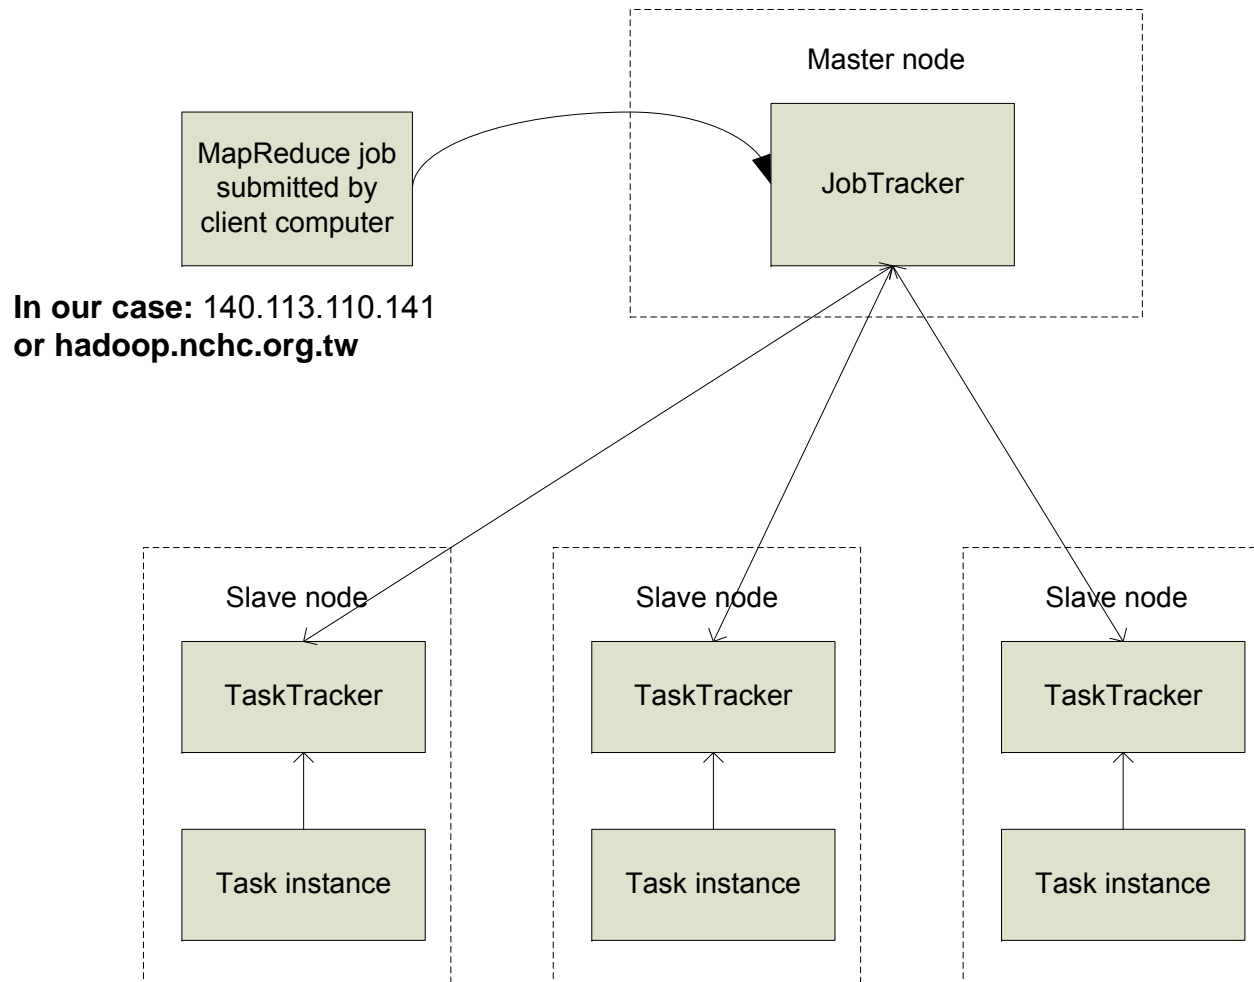
Some MapReduce Terminology

- *Job* – A “full program” - an execution of a Mapper and Reducer across a data set
- *Task* – An execution of a Mapper or a Reducer on a slice of data
 - a.k.a. Task-In-Progress (TIP)
- *Task Attempt* – A particular instance of an attempt to execute a task on a machine

Task Attempts

- A particular task will be attempted at least once, possibly more times if it crashes
 - If the same input causes crashes over and over, that input will eventually be abandoned
- Multiple attempts at one task may occur in parallel with speculative execution turned on
 - Task ID from *TaskInProgress* is not a unique identifier; don't use it that way

MapReduce: High Level



Nodes, Trackers, Tasks

- Master node runs *JobTracker* instance, which accepts *Job* requests from clients
- *TaskTracker* instances run on slave nodes
- TaskTracker forks separate Java processes for task instances

Job Distribution

- MapReduce programs are contained in a Java “jar” file + an XML file containing serialized program configuration options
- Running a MapReduce job places these files into the HDFS and notifies TaskTrackers where to retrieve the relevant program code
- ... Where's the data distribution?

Data Distribution

- Implicit in design of MapReduce!
 - All mappers are equivalent; so map whatever data is local to a particular node in HDFS
- If lots of data does happen to pile up on the same node, nearby nodes will map instead
 - Data transfer is handled implicitly by HDFS

What Happens In Hadoop?

Depth First

Job Launch Process: Client

- Client program creates a *JobConf*
 - Identify classes implementing *Mapper* and *Reducer* interfaces
 - `JobConf.setMapperClass()`, `setReducerClass()`
 - Specify inputs, outputs
 - `FileInputFormat.setInputPath()`,
 - `FileOutputFormat.setOutputPath()`
 - Optionally, other options too:
 - `JobConf.setNumReduceTasks()`,
`JobConf.setOutputFormat()`...

Job Launch Process: *JobClient*

- Pass JobConf to JobClient.runJob() or submitJob()
 - runJob() blocks, submitJob() does not
- *JobClient*:
 - Determines proper division of input into *InputSplits*
 - Sends job data to master *JobTracker* server

Job Launch Process: *JobTracker*

- *JobTracker*:

- Inserts jar and JobConf (serialized to XML) in shared location
- Posts a *JobInProgress* to its run queue

Job Launch Process: *TaskTracker*

- *TaskTrackers* running on slave nodes periodically query *JobTracker* for work
- Retrieve job-specific jar and config
- Launch task in separate instance of Java
 - `main()` is provided by Hadoop

Job Launch Process: Task

- `TaskTracker.Child.main()`:
 - ❑ Sets up the child *TaskInProgress* attempt
 - ❑ Reads XML configuration
 - ❑ Connects back to necessary MapReduce components via RPC
 - ❑ Uses *TaskRunner* to launch user process

Job Launch Process: *TaskRunner*

- *TaskRunner*, *MapTaskRunner*, *MapRunner* work in a daisy-chain to launch your *Mapper*
 - Task knows ahead of time which *InputSplits* it should be mapping
 - Calls *Mapper* once for each record retrieved from the *InputSplit*
- Running the *Reducer* is much the same

Creating the *Mapper*

- You provide the instance of *Mapper*
 - Should extend *MapReduceBase*
- One instance of your Mapper is initialized by the *MapTaskRunner* for a *TaskInProgress*
 - Exists in separate process from all other instances of Mapper – no data sharing!

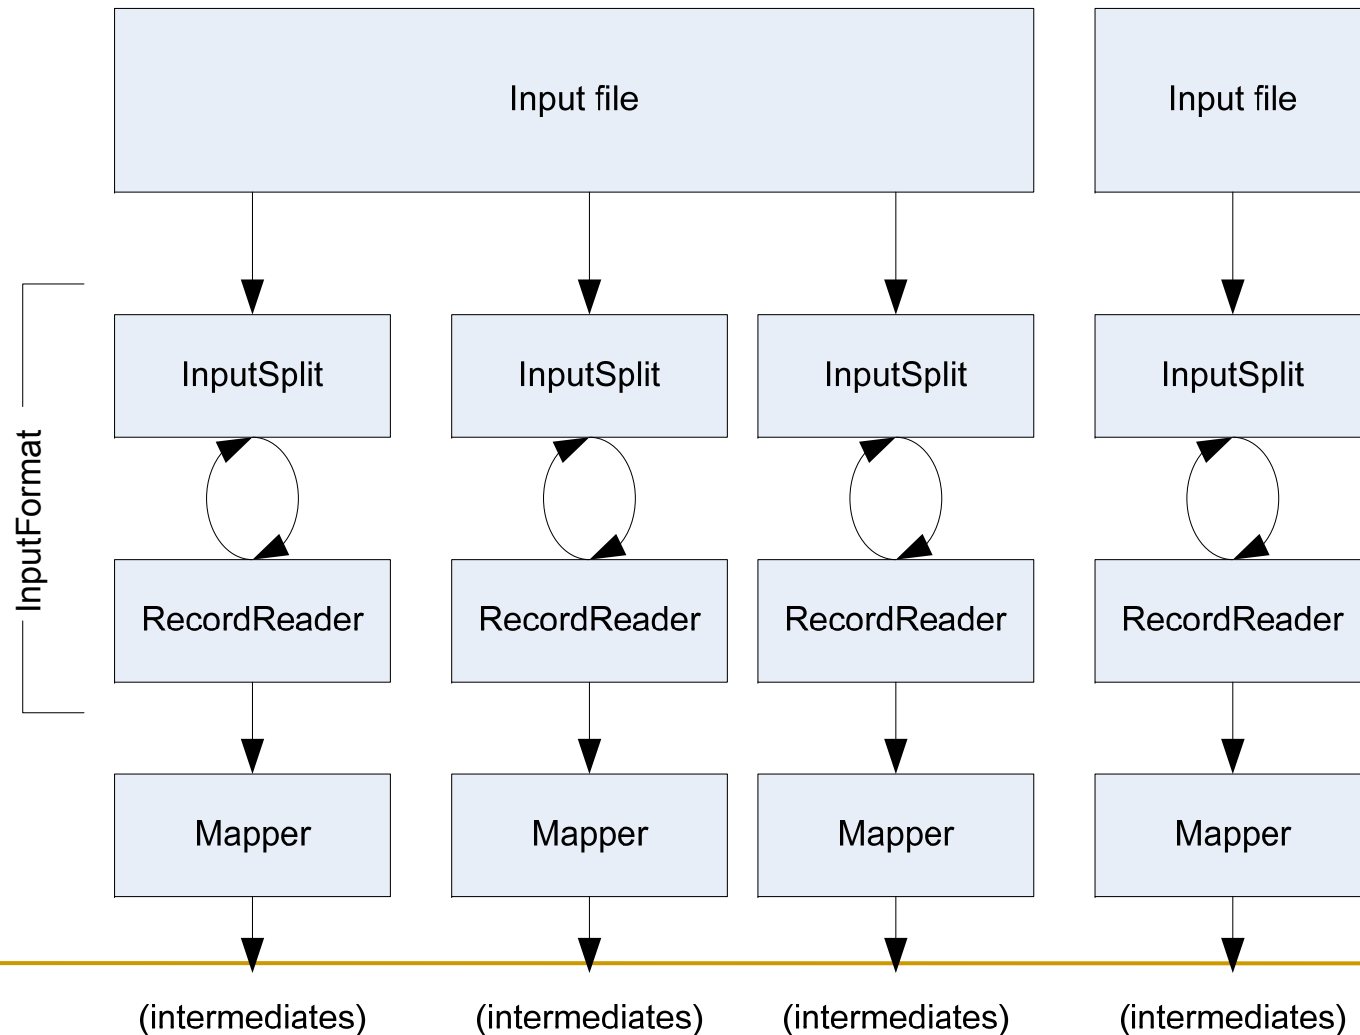
Mapper

- `void map(K1 key,
 V1 value,
 OutputCollector<K2, V2> output,
 Reporter reporter)`
- *K* types implement *WritableComparable*
- *V* types implement *Writable*

What is Writable?

- Hadoop defines its own “box” classes for strings (*Text*), integers (*IntWritable*), etc.
- All values are instances of *Writable*
- All keys are instances of *WritableComparable*

Getting Data To The Mapper



Reading Data

- Data sets are specified by *InputFormats*
 - Defines input data (e.g., a directory)
 - Identifies partitions of the data that form an *InputSplit*
 - Factory for *RecordReader* objects to extract (k, v) records from the input source

FileInputFormat and Friends

- *TextInputFormat* – Treats each ‘\n’-terminated line of a file as a value
- *KeyValueTextInputFormat* – Maps ‘\n’-terminated text lines of “k SEP v”
- *SequenceFileInputFormat* – Binary file of (k, v) pairs with some add’l metadata
- *SequenceFileAsTextInputFormat* – Same, but maps (k.toString(), v.toString())

Filtering File Inputs

- *FileInputFormat* will read all files out of a specified directory and send them to the mapper
- Delegates filtering this file list to a method subclasses may override
 - e.g., Create your own “xyzFileInputFormat” to read *.xyz from directory list

Record Readers

- Each *InputFormat* provides its own *RecordReader* implementation
 - Provides (unused?) capability multiplexing
- *LineRecordReader* – Reads a line from a text file
- *KeyValueRecordReader* – Used by *KeyValueTextInputFormat*

Input Split Size

- *FileInputFormat* will divide large files into chunks
 - Exact size controlled by `mapred.min.split.size`
- RecordReaders receive file, offset, and length of chunk
- Custom *InputFormat* implementations may override split size – e.g., “NeverChunkFile”

Sending Data To Reducers

- Map function receives *OutputCollector* object
 - `OutputCollector.collect()` takes (k, v) elements
- Any (*WritableComparable*, *Writable*) can be used
- By default, mapper output type assumed to be same as reducer output type

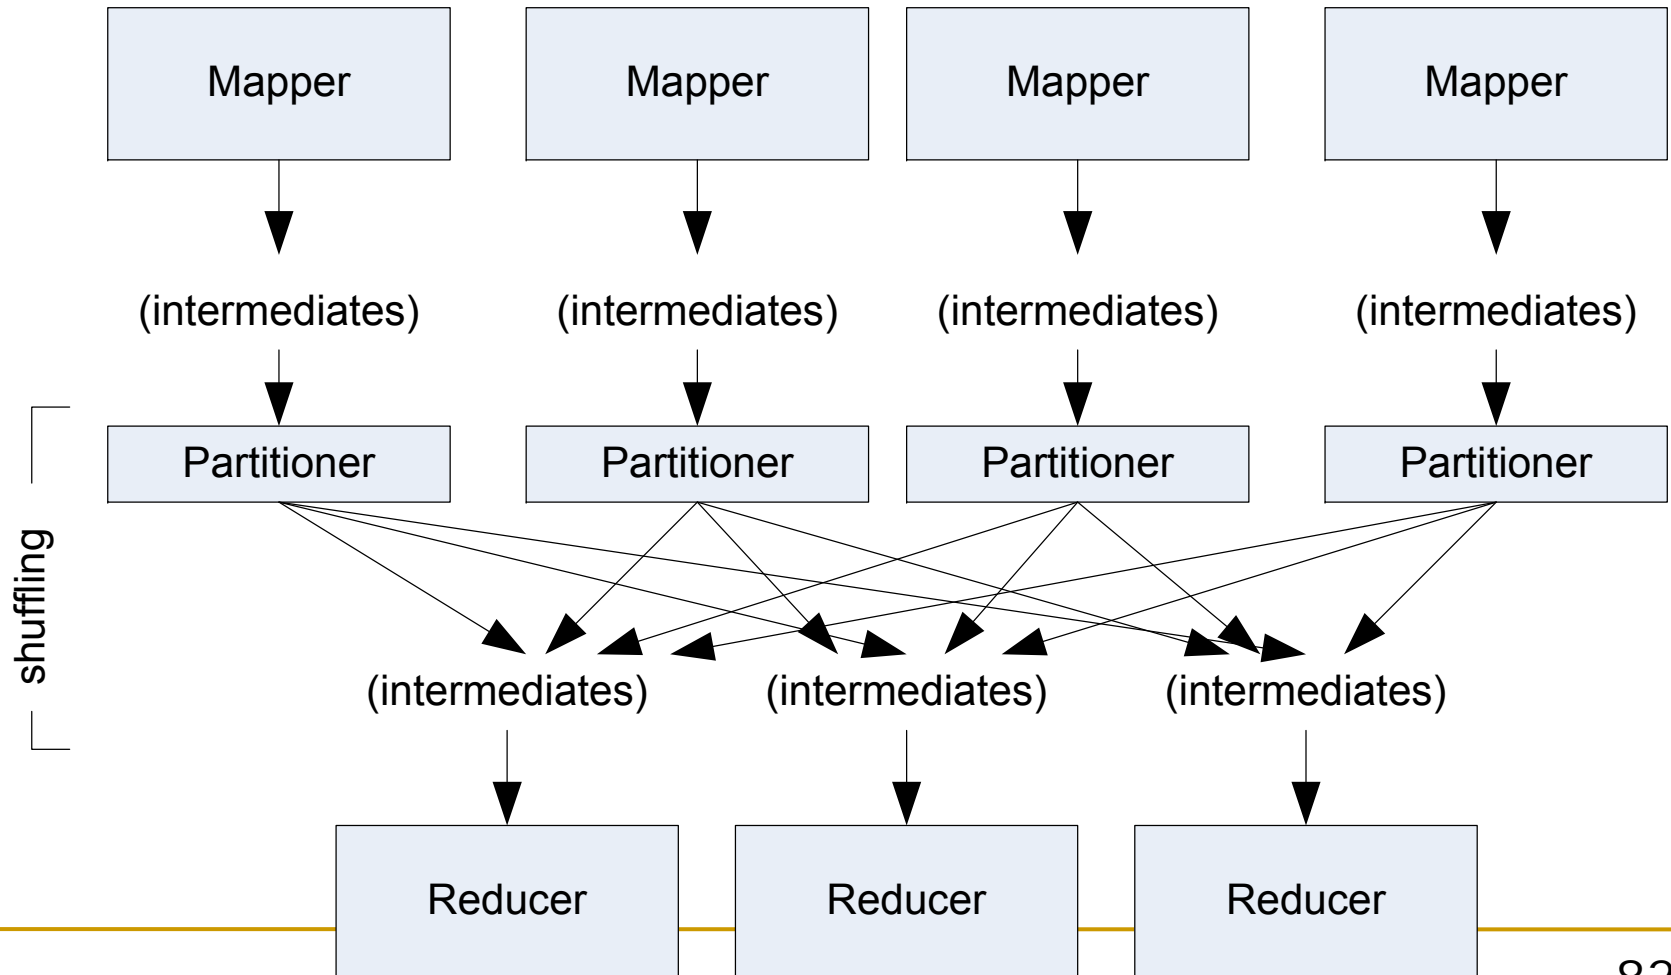
WritableComparator

- Compares WritableComparable data
 - Will call WritableComparable.compare()
 - Can provide fast path for serialized data
- *JobConf.setOutputValueGroupingComparator()*

Sending Data To The Client

- *Reporter* object sent to Mapper allows simple asynchronous feedback
 - `incrCounter(Enum key, long amount)`
 - `setStatus(String msg)`
- Allows self-identification of input
 - `InputSplit getInputSplit()`

Partition And Shuffle



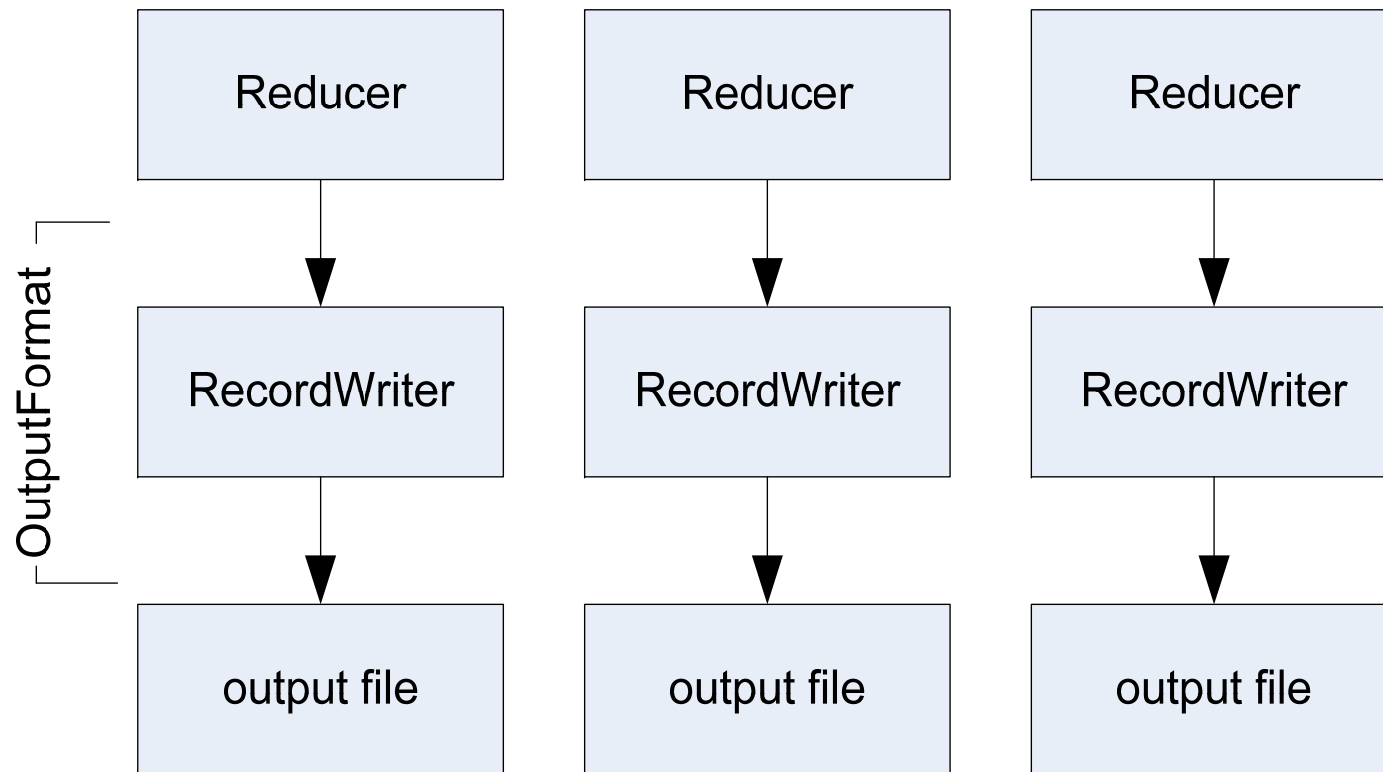
Partitioner

- `int getPartition(key, val, numPartitions)`
 - Outputs the partition number for a given key
 - One partition == values sent to one Reduce task
- *HashPartitioner* used by default
 - Uses `key.hashCode()` to return partition num
- *JobConf* sets *Partitioner* implementation

Reduction

- `reduce(K2 key,`
 `Iterator<V2> values,`
 `OutputCollector<K3, V3> output,`
 `Reporter reporter)`
- Keys & values sent to one partition all go to the same reduce task
- Calls are sorted by key – “earlier” keys are reduced and output before “later” keys

Finally: Writing The Output



OutputFormat

- Analogous to *InputFormat*
- *TextOutputFormat* – Writes “key val\n” strings to output file
- *SequenceFileOutputFormat* – Uses a binary format to pack (k, v) pairs
- *NullOutputFormat* – Discards output
 - Only useful if defining own output methods within `reduce()`

Methods to write MapReduce Jobs

- Typical – usually written in Java
 - MapReduce 2.0 API
 - MapReduce 1.0 API
- Pipes
 - Often used with C++
- Streaming
 - Uses stdin and stdout
 - Can use any language to write Map and Reduce functions
- Abstraction libraries
 - Hive, Pig, etc.
 - Written in a higher level language, generate one or more MapReduce jobs

Example Program - Wordcount

- **map()**
 - Receives a chunk of text
 - Outputs a set of word/count pairs
- **reduce()**
 - Receives a key and all its associated values
 - Outputs the key and the sum of the values

```
package org.myorg;  
import java.io.IOException;  
import java.util.*;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.conf.*;  
import org.apache.hadoop.io.*;  
import org.apache.hadoop.mapred.*;  
import org.apache.hadoop.util.*;
```

```
public class WordCount {
```


Wordcount – main()

```
public static void main(String[] args) throws Exception {  
    JobConf conf = new JobConf(WordCount.class);  
    conf.setJobName("wordcount");  
    conf.setOutputKeyClass(Text.class);  
    conf.setOutputValueClass(IntWritable.class);  
  
    conf.setMapperClass(Map.class);  
    conf.setReducerClass(Reduce.class);  
  
    conf.setInputFormat(TextInputFormat.class);  
    conf.setOutputFormat(TextOutputFormat.class);  
  
    FileInputFormat.setInputPaths(conf, new Path(args[0]));  
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));  
  
    JobClient.runJob(conf);  
}
```

Wordcount – map()

```
public static class Map extends MapReduceBase ... {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value,
                    OutputCollector<Text, IntWritable> output, ...) ... {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            output.collect(word, one);
        }
    }
}
```

Wordcount – reduce()

```
public static class Reduce extends MapReduceBase ... {  
    public void reduce(Text key, Iterator<IntWritable> values,  
                       OutputCollector<Text, IntWritable> output, ...) ... {  
        int sum = 0;  
        while (values.hasNext()) {  
            sum += values.next().get();  
        }  
        output.collect(key, new IntWritable(sum));  
    }  
}  
}
```

Hadoop Pipes: wordcount – main()

- Allows you to create map/reduce jobs in C++

```
#include "hadoop/Pipes.hh"
#include "hadoop/TemplateFactory.hh"
#include "hadoop/StringUtils.hh"

int main(int argc, char *argv[]) {
    return
    HadoopPipes::runTask(HadoopPipes::TemplateFactory<WordCountMap,
                        WordCountReduce>());
}
```

Hadoop Pipes: wordcount – map()

- Allows you to create map/reduce jobs in C++

```
class WordCountMap: public HadoopPipes::Mapper {
public:
    WordCountMap(HadoopPipes::TaskContext& context){}
    void map(HadoopPipes::MapContext& context) {
        std::vector<std::string> words =
            HadoopUtils::splitString(context.getInputValue(), " ");
        for(unsigned int i=0; i < words.size(); ++i) {
            context.emit(words[i], "1");
        }
    }
};
```

Hadoop Pipes: wordcount – reduce()

- Allows you to create map/reduce jobs in C++

```
class WordCountReduce: public HadoopPipes::Reducer {
public:
    WordCountReduce(HadoopPipes::TaskContext& context){}
    void reduce(HadoopPipes::ReduceContext& context) {
        int sum = 0;
        while (context.nextValue()) {
            sum += HadoopUtils::toInt(context.getInputValue());
        }
        context.emit(context.getInputKey(), HadoopUtils::toString(sum));
    }
};
```

Running with Hadoop Pipes

- Allows you to create map/reduce jobs in C++

```
g++ wordcount.cpp -o wordcount -lhadooputils \  
-lhadooppipes -lpthread
```

```
hadoop pipes -D hadoop.pipes.java.recordreader=true \  
-D hadoop.pipes.java.recordwriter=true \  
-program examples/bin/wordcount \  
-input examples/input -output examples/output
```

Hadoop Streaming

- Allows you to create and run map/reduce jobs with any executable
- Similar to unix pipes, e.g.:
 - ❑ format is: Input | Mapper | Reducer
 - ❑ echo “this sentence has five lines” | cat | wc -w

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/hadoop-streaming.jar \  
    -input myInputDirs \  
    -output myOutputDir \  
    -mapper /bin/cat \  
    -reducer /bin/wc
```


Hadoop Streaming

- Mapper and Reducer receive data from stdin and output to stdout
- Hadoop takes care of the transmission of data between the map/reduce tasks
 - It is still the programmer's responsibility to set the correct key/value
 - Default format: "key \t value\n"
- Let's look at a C++/Python example of a MapReduce word count program...

wordcount_map.cpp

```
#include<iostream>
using namespace std;

int main()
{
    string word;
    while (cin >> word)
    {
        cout << word << "\t" << "1" << endl;
    }
    return 0;
}
```

Streaming_Mapper.py

```
# read in one line of input at a time from stdin
for line in sys.stdin:
    line = line.strip()          # string
    words = line.split()        # list of strings

    # write data on stdout
    for word in words:
        print '%s\t%i' % (word, 1)
```

Hadoop Streaming

- What are we outputting?
 - Example output: “the 1”
 - By default, “the” is the key, and “1” is the value
- Hadoop Streaming handles delivering this key/value pair to a Reducer
 - Able to send similar keys to the same Reducer or to an intermediary Combiner

wordcount_reduce.cpp

```
#include<vector>
#include<map>
#include<string>
#include<iostream>
using namespace std;
int main() {
    string key;
    string value;
    map<string,int> mapTemp;
    while(cin >> key >> value) {
        if(mapTemp.find(key) !=
            mapTemp.end())
            mapTemp[key] += 1;
        else
            mapTemp[key] = 1;
    }

    map<string,int>::iterator it = mapTemp.begin();
    for(; it != mapTemp.end(); ++it) {
        cout << it->first << "\t" << it->second << endl;
    }

    return 0;
}
```

Streaming_Reducer.py

```
wordcount = { }          # empty dictionary
# read in one line of input at a time from stdin
for line in sys.stdin:
    line = line.strip()    # string
    key,value = line.split()
    wordcount[key] = wordcount.get(key, 0) + value

# write data on stdout
for word, count in sorted(wordcount.items()):
    print '%s\t%i' % (word, count)
```

Hadoop Streaming Gotcha

- Streaming Reducer receives single lines (which are key/value pairs) from stdin
 - Regular Reducer receives a collection of all the values for a particular key
 - It is still the case that all the values for a particular key will go to a single Reducer

Using Hadoop Distributed File System (HDFS)

- Can access HDFS through various shell commands (see Further Resources slide for link to documentation)
 - ❑ `hadoop -put <localsrc> ... <dst>`
 - ❑ `hadoop -get <src> <localdst>`
 - ❑ `hadoop -ls`
 - ❑ `hadoop -rm file`

Configuring Number of Tasks

- Normal method
 - ❑ `jobConf.setNumMapTasks(400)`
 - ❑ `jobConf.setNumReduceTasks(4)`
- Hadoop Streaming method
 - ❑ `-jobconf mapred.map.tasks=400`
 - ❑ `-jobconf mapred.reduce.tasks=4`
- Note: # of map tasks is only a hint to the framework. Actual number depends on the number of InputSplits generated

Running a Hadoop Job

- Place input file into HDFS:
 - `hadoop fs -put ./input-file input-file`
- Run either normal or streaming version:
 - `hadoop jar Wordcount.jar org.myorg.Wordcount input-file output-file`
 - `hadoop jar hadoop-streaming.jar \`
 - `-input input-file \`
 - `-output output-file \`
 - `-file Streaming_Mapper.py \`
 - `-mapper python Streaming_Mapper.py \`
 - `-file Streaming_Reducer.py \`
 - `-reducer python Streaming_Reducer.py \`

Output Parsing

- Output of the reduce tasks must be retrieved:
 - `hadoop fs -get output-file hadoop-output`
- This creates a directory of output files, 1 per reduce task
 - Output files numbered `part-00000`, `part-00001`, etc.
- Sample output of Wordcount
 - `head -n5 part-00000`

"tis	1
"come	2
"coming	1
"edwin	1
"found	1

Extra Output

- The stdout/stderr streams of Hadoop itself will be stored in an output file (whichever one is named in the startup script)
 - `#$ -o output.$job_id`

STARTUP_MSG: Starting NameNode

STARTUP_MSG: host = svc-3024-8-10.rc.usf.edu/10.250.4.205

...

11/03/02 18:28:47 INFO mapred.FileInputFormat: Total input paths to process : 1

11/03/02 18:28:47 INFO mapred.JobClient: Running job: job_local_0001

...

11/03/02 18:28:48 INFO mapred.MapTask: numReduceTasks: 1

...

11/03/02 18:28:48 INFO mapred.TaskRunner: Task 'attempt_local_0001_m_000000_0' done.

11/03/02 18:28:48 INFO mapred.Merger: Merging 1 sorted segments

11/03/02 18:28:48 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 43927 bytes

11/03/02 18:28:48 INFO mapred.JobClient: map 100% reduce 0%

...

11/03/02 18:28:49 INFO mapred.TaskRunner: Task 'attempt_local_0001_r_000000_0' done.

11/03/02 18:28:49 INFO mapred.JobClient: Job complete: job_local_0001

Further Resources

- Hadoop Tutorial:
<http://developer.yahoo.com/hadoop/tutorial/module1.html>
- Hadoop Streaming:
<http://hadoop.apache.org/common/docs/r0.15.2/streaming.html>
- Hadoop API: <http://hadoop.apache.org/common/docs/current/api>
- HDFS Commands Reference:
http://hadoop.apache.org/hdfs/docs/current/file_system_shell.html

Public Hadoop Cluster

■ <http://hadoop.nchc.org.tw/>

NAR Labs 國家實驗研究院
國家高速網路與計算中心

Hadoop 網頁介面登入

帳號： 密碼：

[新增帳號](#) [忘記密碼](#) [操作問題回報](#)

- 歡迎至 [forum.hadoop.tw](#) 或 [臉書粉絲團](#) 進行討論
- 歡迎加入 [公告群組](#)，以利接收即時公告事宜
- 初學者請參閱[線上教學](#)：[Hadoop 觀念篇](#)，實作步驟請參閱以下三個連結：
[[帳號申請](#) | [HDFS 練習](#) | [MapReduce 練習](#)]
- 本叢集所採用的 Hadoop 版本是 0.20.1，寫程式時請參考 [javahadoop](#)
- 倘若貴單位不允許 SSH 連線，或有 PROXY 限制，請改連 <https://hadoop.nchc.org.tw> 並採用系統帳號密碼(hXXXX)登入。至於 HDFS 與 MapReduce 也請點選對應連結存取。
- R 軟體使用者，可用 <http://hadoop.nchc.org.tw/rstudio> 介面(帳號為 hXXXX)

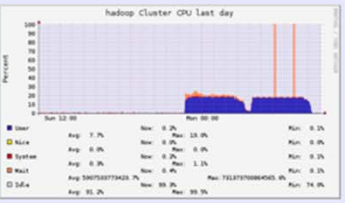
本叢集純為實驗用途，無法保證7x24服務品質
重要數據資料請務必另行備份，謝謝！
重要運算工作亦請考慮使用其他付費平台

家目錄空間吃緊中，請盡量上傳至HDFS後，
清除家目錄檔案，謝謝！

若有公用資料集運算需求，請置於
<hdfs://hadoop.nchc.org.tw/pub>，謝謝！

註冊人數：4294 / 4999 人

[MapReduce 狀態](#) | [HDFS 狀態](#)
過去 24 小時 CPU 負載 - [查詢完整系統負載](#)：



none

Completed Jobs

Jobid	Priority	User	Name	Map % Complete	Map Total	Map Compl
job_201309060021_6629	NORMAL	h4104	wordcount	100.00%	1	1
job_201309060021_6630	NORMAL	h4107	trust	100.00%	1	1
job_201309060021_6633	NORMAL	h4107	trust	100.00%	1	1
job_201309060021_6641	NORMAL	h4105	wordcount	100.00%	1	1
job_201309060021_6643	NORMAL	h4105	wordcount	100.00%	1	1
job_201309060021_6646	NORMAL	h4105	trust	100.00%	1	1
job_201309060021_6647	NORMAL	h4261	wordcount	100.00%	2	2
job_201309060021_6648	NORMAL	h4105	trust	100.00%	1	1
job_201309060021_6649	NORMAL	h4105	trust	100.00%	1	1
job_201309060021_6650	NORMAL	h4105	trust	100.00%	1	1
job_201309060021_6651	NORMAL	h4214	wordcount	100.00%	23	23

NameNode 'hadoop.nchc.org.tw:8020'

Started:	Fri Sep 06 14:25:29 CST 2013
Version:	0.20.1+169.113, r6c765a47a9291470d3d8814c98155115d109d715
Compiled:	Sun Sep 12 05:39:27 UTC 2010 by root
Upgrades:	There are no upgrades in progress.

[Browse the filesystem](#)
[NameNode Logs](#)

Cluster Summary



Reference

- Running Hadoop on Ubuntu Linux (Single-Node Cluster)
 - ◆ <http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-single-node-cluster/>



Conclusions

- MapReduce proven to be useful abstraction
- Greatly simplifies large-scale computations
- Fun to use:
 - ✦ Focus on problem
 - ✦ Let library deal with messy details

